

# 1 Dynamic Partial Order Reduction for Checking Correctness 2 Against Weak Isolation Levels

5 ANONYMOUS AUTHOR(S)

7 Modern applications, such as social networking systems and e-commerce platforms are centered around  
8 using large-scale databases for storing and retrieving data. Accesses to the database are typically enclosed in  
9 transactions that allow computations on shared data to be isolated from other concurrent computations and  
10 resilient to failures. Modern databases trade off isolation for performance. The weaker the isolation level, the  
11 more behaviors a database is allowed to exhibit and it is up to the developer to ensure that their application  
12 can tolerate those behaviors.

13 In this work, we propose a stateless model checking algorithm for studying correctness of such applications  
14 that relies on dynamic partial order reduction. This algorithm works for a number of widely-used weak  
15 isolation levels, including Causal Consistency, Read Committed, and Read Atomic. We show that it is complete,  
16 sound and optimal, and runs with linear memory consumption in all cases. We report on an implementation  
17 of this algorithm in the context of Java Pathfinder applied to a number of challenging applications drawn  
18 from the literature of distributed systems and databases.

## 20 1 INTRODUCTION

21 Programming paradigm is in constant evolution, sequential programs tend to easily become obsolete  
22 because of its slow performance and even concurrent programs can also be inefficient when the  
23 memory requirements increase. The current state-of-the-art tries to overcome those problems by  
24 developing parallel programs along with distributed storage systems. However, not every type  
25 of application has the same data reliability requirements and therefore developers may want to  
26 relax the *isolation level*, i.e. the restrictions imposed to the information stored for guaranteeing  
27 consistency, from the database in order to increase performance.

28 Allowing multiple behaviors in these contexts hinder the already difficult task of verifying con-  
29 current programs. Studying every alternative is something unrealistic, as the number of possible  
30 scenarios grows exponentially with the length of the programs. In general, formal methods such as  
31 cite examples are a reasonable approach as they provide certificate of correctness and explainability  
32 of the bugs otherwise. Among them, *stateless model checking* (SMC) and *dynamic partial order*  
33 *reduction* (DPOR) cite papers stand out as the most promising techniques for verifying current  
34 programs during the recent years cite papers.

35 On one hand, for a given length-bounded program, SMC explores systematically every possible  
36 execution without storing at any point the set of already visited ones. On the other hand, DPOR  
37 resumes every possible behavior in a more succinct way, reducing the number of executions  
38 that have to be explored for covering those behaviors. Henceforth, combining both techniques to  
39 obtaining sound, complete and efficient algorithms has been one of the aimed goals in this field  
40 and it has been successfully done for concurrent programs with shared memory cite papers.

41 Despite their popularity, there is no application of such techniques in parallel programming with  
42 distributed memory's literature so far, hence the relevance of filling this gap. Nevertheless, part  
43 of the path this paper wants to create is already explored, as shared memory models are not that

44 2018. 2475-1421/2018/1-ART1 \$15.00

45 https://doi.org/

46

```

50            $x \in \text{Vars}$     $a \in \text{LVars}$ 
51
52
53      $\text{Prog} ::= \text{Sess} \mid \text{Sess} \parallel \text{Prog}$             $\text{Body} ::= \text{Instr} \mid \text{Instr}; \text{Body}$ 
54      $\text{Sess} ::= \text{Trans} \mid \text{Trans}; \text{Sess}$             $\text{Instr} ::= \text{InstrDB} \mid a := e \mid \text{if}(\phi(\vec{a}))\{\text{Instr}\}$ 
55      $\text{Trans} ::= \text{begin}; \text{Body}; \text{commit}$             $\text{InstrDB} ::= a := \text{read}(x) \mid \text{write}(x, a)$ 
56
57

```

Fig. 1. Program syntax. The set of global variables is denoted by  $\text{Vars}$  while  $\text{LVars}$  denotes the set of local variables. We use  $\phi$  to denote Boolean expressions, and  $e$  to denote expressions over local variables interpreted as values. We use  $\vec{a}$  to denote vectors of elements.

unrelated with distributed database's. For example, we can mention the relation between *sequential consistency* and *serializability* or *strong release-acquire* and *causal consistency*; where both database isolation level cases are nothing but a generalization of their shared memory counterparts [cite papers](#).

In this paper, we present STMC, a *sound, complete, optimal* DPOR algorithm with *linear memory requirements* that employs SMC techniques for verifying some isolation levels. We describe the models that can guarantee those properties, show that *causal consistency* (CC), *read atomic* (RA) and *read committed* (RC) satisfy them and present an example of why more complex models such as *serializability* (SER) cannot be verified with our algorithm. We also present a formal semantics for STMC and exhibit how it evolves from the base algorithm to its current state; requiring executing transactions in isolation and swapping complete blocks of transactions. In addition, we provide some proofs to help the reader having a better understanding of STMC.

On top of this theoretical development, we also furnish this work with an implementation using Java and several benchmarks that study its re. In a nutshell, our software is an extension of JPF [cite tool](#), a Java-built software analysis framework for Java (parallel) programs. It provides control to DFS traversal of executions, which along its modularity, makes it an ideal tool for developing and extending a database concurrent programs' verifier. In particular, we highlight the easiness for splitting the program memory and database's management and providing an API for writing the programs to analyze.

## 2 TRANSACTIONAL PROGRAMS

### 2.1 Program Syntax

Figure 1 lists the definition of a simple programming language that we use to represent applications running on top of a database. A program is a set of *sessions* running in parallel, each session being composed of a sequence of *transactions*. Each transaction is delimited by begin and commit instructions, and its body contains instructions that access the database and manipulate a set  $\text{LVars}$  of local variables.<sup>1</sup> We use symbols  $a, b$ , etc. to denote elements of  $\text{LVars}$ .

For simplicity, we abstract the database state as a valuation to a set  $\text{Vars}$  of *global variables*<sup>2</sup>. Therefore, the instructions accessing the database correspond to reading the value of a global

<sup>1</sup>For simplicity, we assume that all the transactions in the program commit. Aborted transactions can be ignored when reasoning about safety because their effects should be invisible to other transactions.

<sup>2</sup>In the context of a relational database, global variables correspond to fields/rows of a table while in the context of a key-value store, they correspond to keys.

variable and storing it into a local variable  $a$  ( $a := \text{read}(x)$ ), writing the value of a local variable  $a$  to a global variable  $x$  ( $\text{write}(x, a)$ ), or an assignment to a local variable  $a$  ( $a := e$ ). The set of values of global or local variables is denoted by  $\text{Vals}$ . Assignments to local variables use expressions  $e$  over local variables, which are interpreted as values and whose syntax is left unspecified. Each of these instructions can be guarded by a Boolean condition  $\phi(\vec{a})$  over a set of local variables  $\vec{a}$  (their syntax is not important). Our results assume bounded programs, and therefore, we omit other constructs like while loops.

## 2.2 Isolation Level Definition

We present the axiomatic framework introduced by ? for defining isolation levels in key-value stores.<sup>3</sup> Isolation levels are defined as logical constraints, called *axioms*, over *histories*, which are an abstract representation of the interaction between a program and the database in a concrete execution.

**2.2.1 Histories.** Programs interact with a database by issuing transactions formed of begin, end, read and write instructions. The effect of executing one such instruction is represented using an *event*  $\langle e, \text{type} \rangle$  where  $e$  is an *identifier* and *type* is a *type*. There are four types of events: begin, commit,  $\text{read}(x)$  for reading the global variable  $x$ , and  $\text{write}(x, v)$  for writing value  $v$  to global variable  $x$ . The set of events is denoted by  $\mathcal{E}$ . For a read/write event  $e$ , we use  $\text{var}(e)$  to denote the variable  $x$ .

A *transaction log*  $\langle t, E, \text{po}_t \rangle$  is an identifier  $t$  and a finite set of events  $E$  along with a strict total order  $\text{po}_t$  on  $E$ , called *program order*. The minimal element of  $\text{po}_t$  is a begin event. A transaction log without a commit event is called *pending*. Otherwise, it is called *complete*. If the commit event occurs, then it is maximal in  $\text{po}_t$ . The set  $E$  of events in a transaction log  $t$  is denoted by  $\text{events}(t)$ .

The program order  $\text{po}_t$  represents the order between instructions in the body of a transaction. We assume that each transaction log is well-formed in the sense that if a read of a global variable  $x$  is preceded by a write to  $x$  in  $\text{po}_t$ , then it should return the value written by the last write to  $x$  before the read (w.r.t.  $\text{po}_t$ ). This property is implicit in the definition of every isolation level that we are aware of. For simplicity, we may use the term *transaction* instead of *transaction log*.

The set of  $\text{read}(x)$  events in a transaction log  $t$  that are *not* preceded by a write to  $x$  in  $\text{po}_t$ , for some  $x$ , is denoted by  $\text{reads}(t)$ . As mentioned above, the other read events take their values from writes in the same transaction and their behavior is independent of other transactions. Also, the set of  $\text{write}(x, \_)$  events in  $t$  that are *not* followed by other writes to  $x$  in  $\text{po}_t$ , for some  $x$ , is denoted by  $\text{writes}(t)$ . If a transaction contains multiple writes to the same variable, then only the last one (w.r.t.  $\text{po}_t$ ) can be visible to other transactions (w.r.t. any isolation level that we are aware of). The extension to sets of transaction logs is defined as usual. Also, we say that a transaction log  $t$  *writes*  $x$ , denoted by  $t \text{ writes } x$ , when  $\text{writes}(t)$  contains some  $\text{write}(x, \_)$  event.

A *history* contains a set of transaction logs (with distinct identifiers) ordered by a (partial) *session order*  $\text{so}$  that represents the order between transactions in the same session.<sup>4</sup> It also includes a *write-read* relation (also called *read-from*) that “justifies” read values by associating each read to a transaction that wrote the value returned by the read.

<sup>3</sup>Isolation levels are called consistency models by ?.

<sup>4</sup>In the context of our programming language,  $\text{so}$  would be a union of total orders. This constraint is not important for defining isolation levels.

148 *Definition 2.1.* A history  $\langle T, \text{so}, \text{wr} \rangle$  is a set of transaction logs  $T$  along with a strict partial session  
 149 order  $\text{so}$ , and a write-read relation  $\text{wr} \subseteq T \times \text{reads}(T)$  such that

- 150     • the inverse of  $\text{wr}$  is a total function,
- 151     • if  $(t, e) \in \text{wr}$ , where  $e$  is a  $\text{read}(x)$  event, then  $\text{writes}(t)$  contains some  $\text{write}(x, \_)$  event,  
 152       and
- 153     •  $\text{so} \cup \text{wr}$  is acyclic.

156 We assume that every history includes a distinguished transaction log writing the initial values of  
 157 all global variables. This transaction log precedes all the other transaction logs in  $\text{so}$ . We use  $h, h_1,$   
 158  $h_2, \dots$  to range over histories.  
 159

160 For a variable  $x$ ,  $\text{wr}_x$  denotes the restriction of  $\text{wr}$  to reads of  $x$ , i.e.,  $\text{wr}_x = \text{wr} \cap (T \times$   
 161  $\{e \mid e \text{ is a } \text{read}(x) \text{ event}\})$ . Moreover, we extend the relations  $\text{wr}$  and  $\text{wr}_x$  to pairs of transactions by  
 162  $\langle t_1, t_2 \rangle \in \text{wr}$ , resp.,  $\langle t_1, t_2 \rangle \in \text{wr}_x$ , iff there exists a  $\text{read}(x)$  event  $e$  in  $t_2$  such that  $\langle t_1, e \rangle \in \text{wr}$ , resp.,  
 163  $\langle t_1, e \rangle \in \text{wr}_x$ . We say that the transaction log  $t_1$  is *read* by the transaction log  $t_2$  when  $\langle t_1, t_2 \rangle \in \text{wr}$ .

164 The set of transaction logs  $T$  in a history  $h = \langle T, \text{so}, \text{wr} \rangle$  is denoted by  $\text{tr}(h)$ , and the union of  
 165 events( $t$ ) for  $t \in T$  is denoted by  $\text{events}(h)$ . Given a history  $h$  and an event  $e$  in  $h$ ,  $\text{tr}(h, e)$  is the  
 166 transaction  $t$  in  $h$  that contains  $e$ . Also, we define  $\text{writes}(h) = \bigcup_{t \in \text{tr}(h)} \text{writes}(t)$  and  $\text{reads}(h) =$   
 167  $\bigcup_{t \in \text{tr}(h)} \text{reads}(t)$ .

168 To simplify the notation, we extend  $\text{so}$  and  $\text{wr}$  to pairs of events by  $(e_1, e_2) \in \text{so}$  if  
 169  $(\text{tr}(h, e_1), \text{tr}(h, e_2)) \in \text{so}$  and  $(e_1, e_2) \in \text{wr}$  if  $e_1 \in \text{writes}(\text{tr}(h, e_1))$ ,  $(\text{tr}(h, e_1), e_2) \in \text{wr}$ , and  
 170  $\text{var}(e_1) = \text{var}(e_2)$ . We also define  $\text{po} = \bigcup_{t \in T} \text{po}_t$ .

172 TODO NOT SURE THAT THE FOLLOWING DEFINITION IS NEEDED

174 *Definition 2.2.* Let  $h$  be a history:

- 176     •  $h$  is called *complete* if every transaction is non-pending and *incomplete* otherwise;
- 177     •  $h$  is *executed in isolation* if it contains at most one pending transaction;
- 178     •  $h$  is called *total* if it is complete and contains every transaction  $T \in \mathcal{T}$ .

180 ~~If the event  $e = \text{NEXT}(h)$  is begin, write or end, we will denote by  $h \bullet e$  the history  
 181  $h' = \langle E', \text{so}', \text{wr} \rangle$  where  $E' = \text{events}(h) \cup \{e\}$  and  $\text{so}' = \text{so} \cup \{(e', e) \mid e' \in h \wedge \text{th}(e) = \text{th}(e')\}$ . On  
 182 the other hand, if  $e$  is a read event, we will define the history  $h'_w = h \bullet_w e$  for some write event  
 183  $w \in h$  as  $\langle E', \text{so}', \text{wr} \cup \{(w, r)\} \rangle$ , where  $E'$  and  $\text{so}'$  defined as before.: Not well defined, just cut-  
 184 pasted from below.~~

186 EXTENSIONS,  $\mathcal{H}^<$  AND  $\bullet$  OPERATOR NOT (yet) DEFINED IN THIS SECTION!!!!!!

188 2.2.2 *Axiomatic Framework.* To fully model any behavior of a transactional concurrent program  
 189 we are obliged to formally describe the database section. This notion will be depicted as the concept  
 190 of *model*:

192 *Definition 2.3.* An axiomatic *model*  $\mathcal{M}$  over histories is a collection of rules that enforce a *consistency*  
 193 *criterion* over them. The histories that satisfy those criteria are called  $\mathcal{M}$ -consistent while the rest  
 194 are simply denoted  $\mathcal{M}$ -inconsistent. If there is no ambiguity on the model, we will simply denote  
 195 them consistent or inconsistent.

<pre> 197 198 199 200 201 202 203 204 205 206 207 208 209 210 211 212 213 214 215 216 217 </pre> <p><math>\forall x, \forall t_1, t_2, \forall \alpha. t_1 \neq t_2 \wedge \langle t_1, \alpha \rangle \in \text{wr}_x \wedge t_2 \text{ writes } x \wedge \langle t_2, \alpha \rangle \in \text{wr} \circ \text{po} \Rightarrow \langle t_2, t_1 \rangle \in \text{co}</math></p> <p>(a) Read Committed</p>	<pre> writes a t2 --&gt; t3 t2 --&gt; t1 t1 --&gt; t3 t1 --&gt; t2 t2 --&gt; t1 </pre> <p><math>\forall x, \forall t_1, t_2, \forall t_3. t_1 \neq t_2 \wedge \langle t_1, t_3 \rangle \in \text{wr}_x \wedge t_2 \text{ writes } x \wedge \langle t_2, t_3 \rangle \in \text{so} \cup \text{wr} \Rightarrow \langle t_2, t_1 \rangle \in \text{co}</math></p> <p>(b) Read Atomic</p>	<pre> writes x t2 -.-&gt; t3 t2 -.-&gt; t1 t1 --&gt; t3 t1 --&gt; t2 t2 -.-&gt; t1 </pre> <p><math>\forall x, \forall t_1, t_2, \forall t_3. t_1 \neq t_2 \wedge \langle t_1, t_3 \rangle \in \text{wr}_x \wedge t_2 \text{ writes } x \wedge \langle t_2, t_3 \rangle \in (\text{so} \cup \text{wr})^+ \Rightarrow \langle t_2, t_1 \rangle \in \text{co}</math></p> <p>(c) Causal Consistency</p>
<pre> 208 209 210 211 212 213 214 215 216 217 </pre> <p><math>\forall a, \forall t_1, t_2, \forall t_3. t_1 \neq t_2 \wedge \langle t_1, t_3 \rangle \in \text{wr}[a] \wedge t_2 \text{ writes } a \wedge \langle t_2, t_3 \rangle \in \text{co}^* \circ (\text{wr} \cup \text{so}) \Rightarrow \langle t_2, t_1 \rangle \in \text{co}</math></p> <p>(d) Prefix</p>	<pre> writes a t2 --&gt; t4 t2 --&gt; t1 t1 --&gt; t3 t3 --&gt; t4 t4 --&gt; t3 t3 --&gt; t1 </pre> <p><math>\forall a, \forall t_1, t_2, \forall t_3, t_4, \forall y. t_1 \neq t_2 \wedge \langle t_1, t_3 \rangle \in \text{wr}_x \wedge t_2 \text{ writes } a \wedge t_3 \text{ writes } y \wedge t_4 \text{ writes } y \wedge \langle t_2, t_4 \rangle \in \text{co}^* \wedge \langle t_4, t_3 \rangle \in \text{co} \Rightarrow \langle t_2, t_1 \rangle \in \text{co}</math></p> <p>(e) Conflict</p>	<pre> writes x t2 --&gt; t3 t2 --&gt; t1 t1 --&gt; t3 t3 --&gt; t2 t3 --&gt; t4 t4 --&gt; t3 </pre> <p><math>\forall x, \forall t_1, t_2, \forall t_3. t_1 \neq t_2 \wedge \langle t_1, t_3 \rangle \in \text{wr}_x \wedge t_2 \text{ writes } x \wedge \langle t_2, t_3 \rangle \in \text{co} \Rightarrow \langle t_2, t_1 \rangle \in \text{co}</math></p> <p>(f) Serializability</p>

Fig. 2. Axioms defining isolations levels. The reflexive and transitive, resp., transitive, closure of a relation  $rel$  is denoted by  $rel^*$ , resp.,  $rel^+$ . Also,  $\circ$  denotes the composition of two relations, i.e.,  $rel_1 \circ rel_2 = \{(a, b) | \exists c. (a, c) \in rel_1 \wedge (c, b) \in rel_2\}$ .

In figure 2 it is depicted five axioms which correspond to their homonymous isolation levels: *Read Committed* (RC), *Read Atomic* (RA), *Causal Consistency* (CC) *Prefix Consistency* (PRE) and *Serializability* (SER); along with the conflict axiom. Conflict and Prefix allow us to define *Snapshot Isolation* (SI) as the model where prefix and conflict axioms both hold. We say a history  $h$  satisfies an isolation level  $I$  if there is a total order called *commit order*  $co$  that extend  $\text{so} \cup \text{wr}$  and satisfies its axioms. However, by the definition of RC, RA and CC, it is clear that for every history  $h$  s.t. the relation  $co$  deduced from  $\text{so} \cup \text{wr}$  is acyclic exists a commit order for those isolation levels.

TODo DEFINE  $I$ -consistent histories

### 2.3 Program Semantics

We define a small-step operational semantics for transactional programs, which is parametrized by an isolation level  $I$ . The semantics keeps a history of previously executed database accesses in order to maintain conformance to  $I$ .

For readability, we define a program as a partial function  $P : \text{SessId} \rightarrow \text{Sess}$  that associates session identifiers in  $\text{SessId}$  with concrete code as defined in Figure 1 (i.e., sequences of transactions). Similarly, the session order  $\text{so}$  in a history is defined as a partial function  $\text{so} : \text{SessId} \rightarrow \text{Tlogs}^*$  that associates session identifiers with sequences of transaction logs. Two transaction logs are ordered by  $\text{so}$  if one occurs before the other in some sequence  $\text{so}(j)$  with  $j \in \text{SessId}$ .

Formally, the operational semantics is defined as a transition relation  $\Rightarrow_I$  between *configurations*, which are defined as tuples containing the following:

246	<b>SPAWN</b>	$t \text{ fresh } e \text{ fresh } P(j) = \text{begin;Body;commit;S} \quad \vec{B}(j) = \epsilon$
247		$h, \vec{\gamma}, \vec{B}, P \Rightarrow h \oplus_j \langle t, \{ \langle e, \text{begin} \rangle \}, \emptyset \rangle, \vec{\gamma}[j \mapsto \emptyset], \vec{B}[j \mapsto \text{Body;commit}], P[j \mapsto S]$
248		
249		
250	<b>IF-TRUE</b>	$\psi(\vec{x})[x \mapsto \vec{\gamma}(j)(x) : x \in \vec{x}] \text{ true } \quad \vec{B}(j) = \text{if}(\psi(\vec{x})){\{\text{Instr}\}}; B$
251		$h, \vec{\gamma}, \vec{B}, P \Rightarrow h, \vec{\gamma}, \vec{B}[j \mapsto \text{Instr}; B], P$
252		
253	<b>IF-FALSE</b>	$\psi(\vec{x})[x \mapsto \vec{\gamma}(j)(x) : x \in \vec{x}] \text{ false } \quad \vec{B}(j) = \text{if}(\psi(\vec{x})){\{\text{Instr}\}}; B$
254		$h, \vec{\gamma}, \vec{B}, P \Rightarrow h, \vec{\gamma}, \vec{B}[j \mapsto B], P$
255		
256	<b>LOCAL</b>	$v = \vec{\gamma}(j)(e) \quad \vec{B}(j) = a := e; B$
257		$h, \vec{\gamma}, \vec{B}, P \Rightarrow h, \vec{\gamma}[((j, a) \mapsto v)], \vec{B}[j \mapsto B], P$
258		
259		
260	<b>WRITE</b>	$v = \vec{\gamma}(j)(x) \quad e \text{ fresh } \vec{B}(j) = \text{write}(x, a); B \quad h \oplus_j \langle e, \text{write}(x, v) \rangle \text{ satisfies } I$
261		$h, \vec{\gamma}, \vec{B}, P \Rightarrow h \oplus_j \langle e, \text{write}(x, v) \rangle, \vec{\gamma}[j \mapsto B], P$
262		
263		
264	<b>READ-LOCAL</b>	$\text{writes}(last(h, j)) \text{ contains a write}(x, v) \text{ event } \quad e \text{ fresh } \quad \vec{B}(j) = a := \text{read}(x); B$
265		$h, \vec{\gamma}, \vec{B}, P \Rightarrow h \oplus_j \langle e, \text{read}(x) \rangle, \vec{\gamma}[(j, a) \mapsto v], \vec{B}[j \mapsto B], P$
266		
267		
268	<b>READ-EXTERN</b>	$\text{writes}(last(h, j)) \text{ does not contain a write}(x, v) \text{ event } \quad e \text{ fresh } \quad \vec{B}(j) = a := \text{read}(x); B$
269		$h = (T, \text{so}, \text{wr}) \quad t = last(h, j) \quad \text{write}(x, v) \in \text{writes}(t') \text{ with } t' \in \text{compTrans}(h) \text{ and } t \neq t'$
270		$h' = (h \oplus_j \langle e, \text{read}(x) \rangle) \oplus \text{wr}(t', e) \quad h' \text{ satisfies } I$
271		
272		$h, \vec{\gamma}, \vec{B}, P \Rightarrow h', \vec{\gamma}[(j, a) \mapsto v], \vec{B}[j \mapsto B], P$
273		
274	<b>COMMIT</b>	$e \text{ fresh } \quad \vec{B}(j) = \text{commit}$
275		$h, \vec{\gamma}, \vec{B}, P \Rightarrow h \oplus_j \langle e, \text{commit} \rangle, \vec{\gamma}, \vec{B}[j \mapsto \epsilon], P$
276		
277		

Fig. 3. An operational semantics for transactional programs. Above,  $last(h, j)$  denotes the last transaction log in the session order  $\text{so}(j)$  of  $h$ , and  $\text{compTrans}(h)$  denotes the set of transaction logs in  $h$  that are complete.

- history  $h$  storing the events generated by database accesses executed in the past,
- a valuation map  $\vec{\gamma}$  that records local variable values in the current transaction of each session ( $\vec{\gamma}$  associates identifiers of sessions that have live transactions with valuations of local variables),
- a map  $\vec{B}$  that stores the code of each live transaction (associating session identifiers with code), and
- sessions/transactions  $P$  that remain to be executed from the original program.

Before presenting the definition of  $\Rightarrow_I$ , we introduce some notation. Let  $h$  be a history that contains a representation of  $\text{so}$  as above. We use  $h \oplus_j \langle t, E, \text{po} \rangle$  to denote a history where  $\langle t, E, \text{po} \rangle$  is appended to  $\text{so}(j)$ . Also, for a event  $e$ ,  $h \oplus_j e$  is the history obtained from  $h$  by adding  $e$  to the

295 last transaction log in  $\text{so}(j)$  and as a last event in the program order of this log (i.e., if  $\text{so}(j) = \sigma; \langle t, E, \text{po} \rangle$ , then the session order  $\text{so}'$  of  $h \oplus_j e$  is defined by  $\text{so}'(k) = \text{so}(k)$  for all  $k \neq j$  and  
 296  $\text{so}(j) = \sigma; \langle t, E \cup \{e\}, \text{po} \cup \{(e', e) : e' \in E\} \rangle$ ). Finally, for a history  $h = \langle T, \text{so}, \text{wr} \rangle$ ,  $h \oplus \text{wr}(t, e)$  is  
 297 the history obtained from  $h$  by adding  $(t, e)$  to the write-read relation.

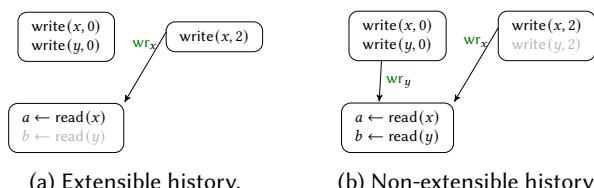
298  
 299 Figure 3 lists the rules defining  $\Rightarrow_I$ . SPAWN starts a new transaction in a session  $j$  provided that  
 300 this session has no live transaction ( $\vec{B}(j) = \epsilon$ ). It adds a transaction log with a single begin event to  
 301 the history and schedules the body of the transaction. IF-TRUE and IF-FALSE check the truth value  
 302 of a Boolean condition of an if conditional. LOCAL models the execution of an assignment to a  
 303 local variable which does not impact the stored history. READ-LOCAL and READ-EXTERN concern  
 304 read instructions. READ-LOCAL handles the case where the read follows a write on the variable  $x$  in  
 305 the same transaction: the read returns the value written by the last write on  $x$  in that transaction.  
 306 Otherwise, READ-EXTERN corresponds to reading a value written in another transaction  $t'$ . The  
 307 transaction  $t'$  is chosen non-deterministically as long as extending the current history with the  
 308 write-read dependency associated to this choice leads to a history that still satisfies  $I$ . READ-EXTERN  
 309 applies only when the executing transaction contains no write on the same variable.  
 310

311 An *initial* configuration for program  $P$  contains the program  $P$  along with a history  $h = \langle \{t_0\}, \emptyset, \emptyset \rangle$ ,  
 312 where  $t_0$  is a transaction log containing only writes that write the initial values of all variables, and  
 313 empty current transaction code ( $B = \epsilon$ ). An execution of a program  $P$  under an isolation level  $I$  is a  
 314 sequence of configurations  $c_0 c_1 \dots c_n$  where  $c_0$  is an initial configuration for  $P$ , and  $c_m \Rightarrow_I c_{m+1}$ ,  
 315 for every  $0 \leq m < n$ . We say that  $c_n$  is  $I$ -reachable from  $c_0$ . The history of such an execution is  
 316 the history  $h$  in the last configuration  $c_n$ . A configuration is called *final* if it contains the empty  
 317 program ( $P = \emptyset$ ). Let  $\text{hist}_I(P)$  denote the set of all histories of an execution of  $P$  under  $I$  that ends  
 318 in a final configuration.  
 319

### 320 3 PREFIX-CLOSED AND CAUSALLY-EXTENSIBLE ISOLATION LEVELS

321 TODO THE INTUITION BEHIND THIS CONDITION (USE THE TEXT THAT FOLLOWS)

322 Besides models presented in figure 2, others isolation levels exists in literature and real life ap-  
 323 plications cite Constantin's papers + Twitter, shoppingcart... However, our algorithm can not be  
 324 analyzed under an arbitrary model. We characterize in this section the ones that can be employed  
 325 by our algorithm.  
 326



350 Fig. 4. Example of a dead-lock after swapping two events.

351 Let's analyze the histories  $h_1$  and  $h_2$  described in figure 4a and 4b respectively under RA; isolation  
 352 level under which both are consistent.  $h_1$  can be extended adding the event  $r_1 = b \leftarrow \text{read}(y)$  and  
 353 the wr-edge  $w_1$  [wr]  $r_1$ , where  $w_1 = \text{write}(x, 0)$ . However, this is not the case of  $h_2$ : the only event  
 354 that could be added in it is  $w_2 = \text{write}(y, 2)$ . If  $w_2$  would be added in  $h_2$ , any relation extending  
 355  $\text{so} \cup \text{wr}$  and satisfying RA would be cyclic, so it wouldn't be a commit order. The essential difference  
 356 between these two histories is the following: in  $h_1$ ,  $t(r)$  is  $\text{so} \cup \text{wr}$ -maximal while in  $h_2$   $t(w)$  is not.  
 357  
 358

344 As real database executions forbid transactions reading from non-committed ones, it is reasonable  
 345 to allow those transactions  $\text{so} \cup \text{wr}$ -maximal to be executed completed without hindering the  
 346 previous committed transactions.

347 For a relation  $R \subseteq A \times A$ , the restriction of  $R$  to  $A' \times A'$ , denoted by  $R \downarrow A' \times A'$ , is defined by  
 348  $\{(a, b) : (a, b) \in R, a, b \in A'\}$ . Also, a set  $A'$  is called  $R$ -downward closed when it contains  $a \in A$   
 349 every time it contains some  $b \in A$  with  $(a, b) \in R$ .

350 A *prefix* of a transaction log  $\langle t, E, \text{po}_t \rangle$  is a transaction log  $\langle t, E', \text{po}_t \downarrow E' \times E' \rangle$  such that  
 351  $E'$  is  $\text{po}_t$ -downward closed. A *prefix* of a history  $h = \langle T, \text{so}, \text{wr} \rangle$  is a history  $h' =$   
 352  $\langle T', \text{so} \downarrow T' \times T', \text{wr} \downarrow T' \times T' \rangle$  such that every transaction log in  $T'$  is a prefix of a different trans-  
 353 action log in  $T$  but carrying the same id, and events( $h'$ )  $\subseteq$  events( $h$ ) is  $(\text{po} \cup \text{so} \cup \text{wr})^*$ -downward  
 354 closed. TODO GIVE AN EXAMPLE OF A PREFIX ON THE FIGURE ABOVE.

355  
 356  
 357 Definition 3.1. An isolation level  $I$  is called *prefix-closed* when every prefix of an  $I$ -consistent history  
 358 is also  $I$ -consistent.

359  
 360 THEOREM 3.2. Every model depicted in figure 2 is prefix closed.

361 PROOF. Let  $h$  be a consistent history. As any  $\text{so} \cup \text{wr}$ -prefix-closed sub-history  $h'$  of  $h$  is a sub-graph  
 362 of it, and there is a commit order  $\text{co}$  for  $h$ , it suffices to restrict  $\text{co}$  to  $h'$  for obtaining a commit order  
 363 for  $h'$ .  $\square$

364  
 365 Let  $h = \langle T, \text{so}, \text{wr} \rangle$  be a history. A transaction  $t$  is called  $(\text{so} \cup \text{wr})^*$ -maximal in  $h$  if  $h$  does not  
 366 contain any transaction  $t'$  such that  $(t, t') \in (\text{so} \cup \text{wr})^*$ . We define a *causal extension* of a pending  
 367 transaction  $t$  in  $h$  with an event  $e$  as follows:

- 368  
 369  
 370   •  $e$  is added to  $t$  as a maximal element of  $\text{po}_t$ ,  
 371   • if  $e$  is a read event, then  $\text{wr}$  is extended with some tuple  $(t', e)$  such that  $t' [\text{so} \cup \text{wr}]^* t$   
 372   • the other elements of  $h$  remain unchanged.

373  
 374  
 375 TODO NOTE THAT A HISTORY MAY HAVE MULTIPLE CAUSAL EXTENSIONS WITH A READ.  
 376 NOT FOR THE OTHER TYPES OF EVENTS. GIVE EXAMPLES.

377  
 378 Definition 3.3. An isolation level  $I$  is called *causally-extensible* if for every  $I$ -consistent history  
 379  $h$ , every  $(\text{so} \cup \text{wr})^*$ -maximal pending transaction  $t$  in  $h$ , and every event  $e$ , there exists a causal  
 380 extension of  $t$  with  $e$  that is  $I$ -consistent.

381  
 382 THEOREM 3.4. Causal Consistency (CC), Read Atomic (RA) and Read Committed (RC) are causally-  
 383 extensible.

384  
 385 PROOF. Let  $I$  an isolation level in  $\{\text{CC}, \text{RA}, \text{RC}\}$ ,  $h$  a non-total consistent history and let  $e$  a  $\text{so} \cup \text{wr}$ -  
 386 maximal event. If  $e$  is a begin event,  $h \bullet e$  is consistent as if there exists a commit order  $\text{co}$  for  $h$ ,  
 387 the relation  $\text{co}' = \text{co} \cup \{\langle T, t(e) \rangle, T \in h\}$  is a commit order to  $h'$ . Moreover, if  $e$  is either a write  
 388 or an end event,  $h \bullet e$  is edge-wise identical to  $h$ , so the commit order for  $h$  is also a valid commit  
 389 order for  $h \bullet e$ . Therefore, let  $e$  a  $\text{so} \cup \text{wr}$ -maximal read event that reads variable  $x$  and let us find a  
 390 write event  $w$  s.t.  $t(w) [\text{so} \cup \text{wr}]^* t(r)$  and that  $h'_w = h \bullet_w r$  is consistent.

393 For doing so, we will do an induction on the number of **co** cycles  $h'_w$  has, for some event  $w$  s.t.  
 394  $t(w) [\text{so} \cup \text{wr}]^* t(r)$ ; where **co** Clearly, if  $h'_w$  is acyclic, by theorem **cite theorem h acyclic => exists**  
 395 **a co (another paper, I hope it exists somewhere)**, it is consistent. Hence, let's suppose that if  $h'_w$  has  
 396 at most  $n$  cycles, there exists another write event  $w_n$  s.t.  $r$  causally depends on and  $h'_{w_n} = h \bullet_{w_n} r$  with  
 397 is consistent; and let's analyze if the same property can be deduced for a history  $h'_{w_{n+1}} = h \bullet_{w_{n+1}} r$  with  
 398  $n+1$  cycles. As  $h$  is consistent and  $t(w_{n+1}) [\text{so} \cup \text{wr}]^* t(r)$ , if there were a cycle, it would be due to a  
 399 transaction  $T$  such that writes  $x$ ,  $t(w_{n+1}) [\text{co}]^* T$  and  $\varphi_M(T, e)$ , where  $\varphi_{CC}(T, e) = T [\text{so} \cup \text{wr}]^+ t(e)$ ,  
 400  $\varphi_{RA}(T, e) = T [\text{so} \cup \text{wr}] t(e)$  and  $\varphi_{RC}(T, e) = T [\text{wr} \circ \text{po}] e$ . In particular, for any of the three models,  
 401  $T [\text{so} \cup \text{wr}]^* t(r)$ . Let  $w_n$  a write event in  $T$  that writes  $x$  and  $h_{w_n} = h \bullet_w r$ : if we prove that  
 402 the number of **co**-cycles in  $h_{w_n}$  is strictly smaller than in  $h_{w_{n+1}}$  by induction hypothesis, we can  
 403 conclude the result.

404 Firstly, as  $T [\text{wr}_x] t(r)$ , the cycle that was between  $T$  and  $t(w_{n+1})$  does not exist in  $h_{w_n}$ . Moreover,  
 405 analogously as before, if there is a cycle in  $h_{w_n}$ , it is due to the existence of a transaction  $T'$  s.t.  
 406  $T [\text{co}] T'$  and  $\varphi_M(T', r)$ . Therefore, in  $h_{w_{n+1}}$   $t(w_{n+1}) [\text{co}] T [\text{co}] T'$ , so there is a cycle between in  
 407  $t(w_{n+1})$  and  $T'$ . To sum up, every cycle in  $h_{w_n}$  has a counterpart in  $h_{w_{n+1}}$  and it has one less cycle;  
 408 so the inductive step holds.

409  
 410 Finally, as **init** contains a write event  $w_0$  that writes  $x$ , **init**  $[\text{so} \cup \text{wr}]^* t(r)$  and every history  
 411 has a finite number of transactions, we can deduce from the history  $h'_{w_0} = h \bullet_{w_0} r$  that there is a  
 412 write event  $w$  that  $r$  depends causally on and  $h \bullet_w r$  is consistent.  $\square$

#### 413 THEOREM 3.5. Prefix Consistency (PRE) is maximally-extensible.

414 PROOF. Let  $h$  a non-total PRE-consistent history, **co** a commit order that witness this property and  
 415  $e$  a **so**  $\cup$  **wr**-maximal event. As if  $e$  is the begin of a transaction,  $\text{co}' = \text{co} \cup \{\langle t, \text{tr}(h, e) \rangle, t \in h\}$  is a  
 416 witness of  $h \oplus e$ 's PRE-consistency; let us show that if  $\text{begin}(\text{tr}(h, e)) \in h$ , there is a commit order  
 417  $\text{co}'$  for  $h$ . If that would be the case,  $h \oplus e$  will be consistent with  $\text{co}'$  as witness.

418 Let  $\text{co}' = \{\langle t, t' \rangle \mid t \times t' \in h^2$  s.t.  $t [\text{co}] t' \wedge t \neq \text{tr}(h, e)\} \cup \{\langle t, \text{tr}(h, e) \rangle \mid t \in h\}$ . As  $\text{co}'$  is a total  
 419 order between transactions,  $\text{co}'$  witness  $h$  is PRE-consistent if and only if there is no  $t_1, t_2, t_3, t_4$   
 420 transactions such that  $(t_1, t_3) \in \text{wr}_x$ ,  $t_2$  writes  $x$ ,  $(t_2, t_4) \in \text{co}'^*$  and  $(t_4, t_3) \in \text{so} \cup \text{wr}$  but  $(t_1, t_2) \in \text{co}'$ .  
 421 If  $\text{tr}(h, e)$  is not either  $t_1, t_2, t_3$  or  $t_4$ , as  $\text{co}'$  only modifies the relative order between  $\text{tr}(h, e)$  and  
 422 every other transaction, this situation cannot happen. On one hand, as  $\text{tr}(h, e)$  is **so**  $\cup$  **wr**-maximal,  
 423  $\text{tr}(h, e)$  cannot be  $t_1$  nor  $t_4$ . On the other hand, as  $\text{tr}(h, e)$  is  $\text{co}'$ -maximum, it cannot be also  $t_1$ .  
 424 Finally, if  $t_3 = \text{tr}(h, e)$ , then the exact situation would happen regarding **co**, which is impossible.  
 425 Therefore,  $\text{co}'$  witness  $h$  is PRE-consistent.  $\square$

426 TODO SHOW THAT SERIALIZABILITY IS NOT CAUSALLY-EXTENSIBLE.

## 427 4 SWAPPING-BASED MODEL CHECKING ALGORITHMS

428 We present a class of model checking algorithms for enumerating executions of a given transactional  
 429 program, that we call *swapping-based algorithms*. Section 5 will describe a concrete instance that  
 430 applies to isolation levels that are prefix-closed and causally extensible.

431 These algorithms are defined by the recursive function EXPLORE listed in Algorithm 1. The function  
 432 EXPLORE receives as input a program  $P$ , an *ordered history*  $h_<$ , which is a pair  $(h, <)$  of a history  
 433 and a total order  $<$  on all the events in  $h$ , and a mapping  $\text{locals}$  that associates each event  $e$  in  $h$   
 434 with the valuation of local variables in the transaction of  $e$  ( $\text{tr}(h, e)$ ) just before executing  $e$ . For  
 435 an ordered history  $(h, <)$  with  $h = \langle T, \text{so}, \text{wr} \rangle$ , we assume that  $<$  is consistent with **po**, **so**, and **wr**,

---

**Algorithm 1** EXPLORE algorithm
 

---

```

442 1: function EXPLORE( $P, h_<, \text{locals}$ )
443 2:    $j, e, \gamma \leftarrow \text{NEXT}(P, h_<, \text{locals})$ 
444 3:    $\text{locals}' \leftarrow \text{locals}[e \mapsto \gamma]$ 
445 4:   if  $e = \perp$  and VALID( $h$ ) then
446 5:     output  $h, \text{locals}'$ 
447 6:   else if type( $e$ ) = read then
448 7:     for all  $w \in \text{VALIDWRITES}(h, e)$  do
449 8:        $h'_< \leftarrow h_< \oplus_j e \oplus \text{wr}(w, e)$ 
450 9:       EXPLORE( $P, h'_<, \text{locals}'$ )
451 10:      EXPLORESWAPS( $P, h'_<, \text{locals}'$ )
452 11:   else
453 12:      $h'_< \leftarrow h_< \oplus_j e$ 
454 13:     EXPLORE( $P, h'_<, \text{locals}'$ )
455 14:     EXPLORESWAPS( $P, h'_<, \text{locals}'$ )
456 15:   end if
457 16: end function
458 17: function EXPLORESWAPS( $P, h_<, \text{locals}$ )
459 18:    $l \leftarrow \text{COMPUTEREORDERINGS}(h_<)$ 
460 19:   for all  $(\alpha, \beta) \in l$  do
461 20:     if OPTIMALITY( $h_<, \alpha, \beta$ ) then
462 21:       EXPLORE( $P, \text{SWAP}(h_<, \alpha, \beta, \text{locals})$ )
463 22:     end if
464 23: end function
465
466
467

```

---

i.e.,  $e_1 < e_2$  if  $(\text{tr}(h, e_1), \text{tr}(h, e_2)) \in (\text{so} \cup \text{wr})^+$  or  $(e_1, e_2) \in \text{po}$ . Initially, the ordered history and the mapping locals are empty.

The function EXPLORE calls NEXT to obtain an event representing the next database access in some pending transaction of the input program. A typical implementation of NEXT would choose one of the pending transactions, execute all local instructions until the next database instruction in that transaction (applying the transition rules IF-TRUE, IF-FALSE, and LOCAL from Figure 3) and return the event  $e$  corresponding to that database instruction and the current local state  $\gamma$ . NEXT may also return  $\perp$  if the program finished. If NEXT returns  $\perp$ , then the function VALID can be used to filter executions that satisfy the intended isolation level before outputting the current history and local states.

Otherwise, the event  $e$  is added to the ordered history  $h_<$ . If  $e$  is a read event, then VALIDWRITES computes a set of write events  $w$  in the current history that are valid for  $e$ , i.e., adding the event  $e$  along with the wr dependency  $(w, e)$  leads to a history that still satisfies the intended isolation level. Concerning notations, we extend  $\oplus_j$  to ordered histories in a straightforward manner:  $(h, <) \oplus_j e$  is the ordered history  $(h \oplus_j e, < \cdot e)$  where  $< \cdot e$  means that  $e$  is added as the maximal element of the total order  $<$ . Also,  $h \oplus_j (e, \text{begin})$  is the same as  $h \oplus_j \langle t, \{(e, \text{begin})\}, \emptyset \rangle$  for a fresh transaction id  $t$ . Moreover, we simplify the notation and write  $h \oplus \text{wr}(w, e)$  for  $w$  a write event and  $e$  a read event instead of  $h \oplus \text{wr}(\text{tr}(h, w), e)$ .

Moreover, once an event is added to the current history, the algorithm may explore other histories obtained by re-ordering events in the current one. Such re-orderings are required for completeness.

491 As explained above, new read events can only read from writes executed in the past which limits  
 492 the set of explored histories to the scheduling imposed by NEXT. Without re-orderings, write events  
 493 scheduled later by NEXT cannot be read by read events executed in the past, even-though this may  
 494 be permitted by the intended isolation level.

495 The function EXPLORESWAPS calls COMPUTEORDERINGS to compute pairs of sequences of events  
 496  $\alpha, \beta$  that should be re-ordered;  $\alpha$  and  $\beta$  are *contiguous and disjoint* subsequences of the total order  
 497  $<$ , and  $\alpha$  should end before  $\beta$  (since  $\beta$  will be re-ordered before  $\alpha$ ). Typically,  $\alpha$  should contain  
 498 a read event  $r$  and  $\beta$  a write event  $w$  such that re-ordering the two enables  $r$  to read from  $w$ .  
 499 Avoiding redundancy, i.e., exploring the same history multiple times, may require restricting the  
 500 application of such re-orderings. This is modeled by the Boolean condition called OPTIMALITY. If  
 501 this condition holds, the new explored histories are computed by the function SWAP. This function  
 502 returns local states as well, which are necessary for continuing the exploration. We assume that  
 503 SWAP( $(h, <), \alpha, \beta, \text{locals}$ ) returns pairs  $((h', <'), \text{locals}')$  such that  
 504

- 505 •  $h'$  contains the events in  $\alpha$  and  $\beta$ ,
- 506 •  $h'$  without the events in  $\alpha$  is a prefix of  $h$ , and
- 507 • for every read event  $r$  in  $\alpha$ , if it reads from different writes in  $h$  and  $h'$  (i.e., the writes  $w_1$   
 508 and  $w_2$  associated with  $r$  by the wr relations of  $h$  and  $h'$ , respectively, are different) then it  
 509 is the last event in its transaction log (w.r.t. po).

510 The first condition makes the re-ordering “meaningful”, while the last two conditions ensure that the  
 511 history  $h'$  is feasible (i.e., it can be obtained using the operational semantics defined in Section 2.3)  
 512 and restricts the possible changes of the wr relation to events in  $\alpha$ . Concerning the events in  $\beta$ , they  
 513 imply that  $h'$  contains all their (po  $\cup$  so  $\cup$  wr)\* predecessors. The last condition is required since  
 514 changing the value returned by a read access may disable later events in the same transaction.  
 515

516 A concrete implementation of EXPLORE is called:

- 517 • *I-sound* if it outputs only histories in  $\text{hist}_I(P)$  for every program  $P$ ,
- 518 • *I-complete* if it outputs every history in  $\text{hist}_I(P)$  for every program  $P$ ,
- 519 • *optimal* if it does not output the same history twice,
- 520 • *strongly optimal* if it is optimal and never engages in fruitless explorations, i.e., EXPLORE is  
 521 never called (recursively) on a history  $h$  that does not satisfy  $I$ , and every call to EXPLORE  
 522 results in an output or a recursive call to EXPLORE.

## 523 5 SWAPPING-BASED MODEL CHECKING FOR PREFIX-CLOSED AND 524 CAUSALLY-EXTENSIBLE ISOLATION LEVELS

525 We present a concrete implementation of EXPLORE that is *I*-sound, *I*-complete, and strongly opti-  
 526 mal for any isolation level  $I$  that is prefix-closed and causally-extensible. Moreover, the space  
 527 complexity of this algorithm is polynomial in the size of the program. An important invariant  
 528 of this implementation is that it explores histories with *at most one* pending transaction and this  
 529 transaction is maximal w.r.t. session order. This invariant is used to avoid fruitless explorations:  
 530 since  $I$  is assumed to be causally-extensible, there always exists an extension of the current history  
 531 with one more event that continues to satisfy  $I$ . We prove the correctness of the algorithm in  
 532 Section 8.

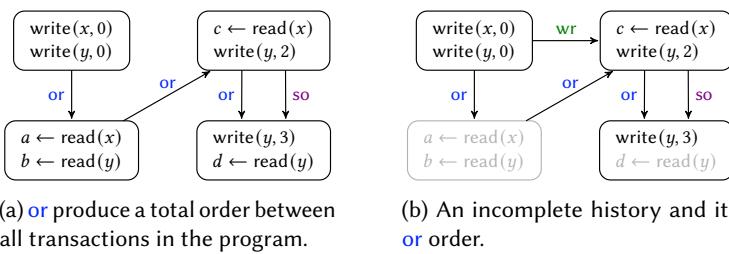
540    **5.1 Extending Histories According to An Oracle Order**

541    The function NEXT generating events that represent database accesses is parametrized by an  
 542    *arbitrary but fixed* order between the transactions in the program called *oracle order*. This order,  
 543    denoted as  $<_{\text{or}}$  and trivially extensible to events in a history, has to respect the order between  
 544    transactions in the same session of the program.

545    NEXT returns a new event of the transaction that is not already completed and that is *minimal*  
 546    according to  $<_{\text{or}}$ . In more detail, if  $j, e, \gamma$  is the output of  $\text{NEXT}(P, h_{<}, \text{locals})$ , then either:

- 548    • the last transaction log  $t$  of session  $j$  (w.r.t. *so*) in  $h$  is pending, and  $t$  is the smallest among  
 549    pending transaction logs in  $h$  w.r.t.  $<_{\text{or}}$
- 550    •  $h$  contains no pending transaction logs and the next transaction of sessions  $j$  is the smallest  
 551    among not yet started transactions in the program.

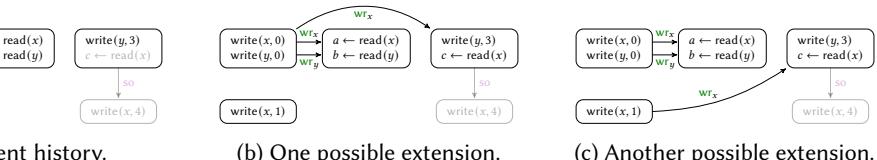
553    This implementation of NEXT is deterministic and it prioritizes the completion of pending transac-  
 554    tions. The latter is useful to maintain the invariant that any history explored by the algorithm has  
 555    at most one pending transaction. Preserving this invariant requires that the histories given as input  
 556    to NEXT also have at most one pending transaction. This is discussed further when explaining the  
 557    process of re-ordering events in Section 5.2.



568    Fig. 5. Some possible oracle order between transactions. TODO MAY NEED TO EXPLAIN THE CONVEN-  
 569    TIONS IN THE PICTURE.

570  
 571    TODO: THIS SHOULD BE AN EXAMPLE OF A STEP OF NEXT. ADD TO THE ABOVE EXAMPLE  
 572    SOME LOCAL INSTRUCTION BEFORE THE READ(Y), AND TALK ABOUT LOCAL STATES AS  
 573    WELL. ALSO, ON THE LEFT, PUT THE PROGRAM AND NOT A HISTORY (PROGRAM WRITTEN  
 574    WITH THE SYNTAX PRESENTED AT THE BEGINNING OF THE PAPER).

576    For example, given the history  $h$  in Figure 5b, the function NEXT would return the event  $d ← \text{read}(y)$   
 577    instead of  $a ← \text{read}(x)$ ; as the forth transaction is pending in  $h$ .



586    Fig. 6. Extensions of a history by adding a read event. TODO MAY NEED TO EXPLAIN THE CONVENTIONS  
 587    IN THE PICTURE.

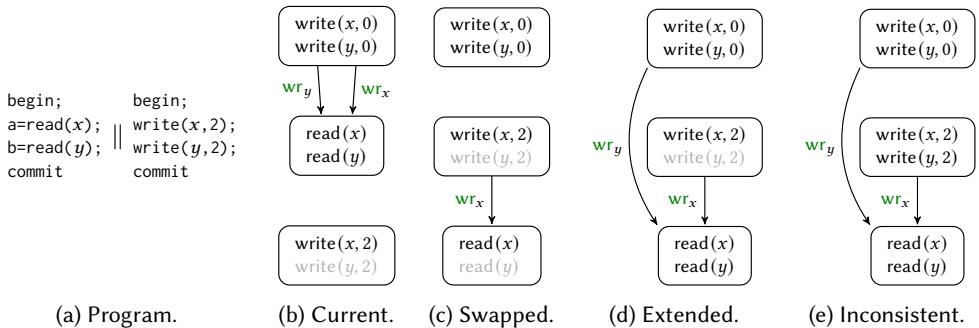
589 TODO REDO THE ABOVE FIGURE IN THE SPIRIT OF FIGURE 7. MODIFY THE EXAMPLE SO  
 590 ONE OF THE TWO CHOICES DOES NOT SATISFY SOME ISOLATION LEVEL, TO EXEMPLIFY  
 591 THE USE OF VALIDWRITES.

592 If the event returned by NEXT is not a read event, then it is simply added to the current history, as  
 593 the maximal element of the order  $<$  (cf. the definition of  $\oplus_j$  on ordered histories). If it is a read event,  
 594 then adding this event may result in multiple histories depending on the associated  $wr$  dependency.  
 595 For example, in Figure 6, extending the history in Figure 6a with the read( $x$ ) event could result  
 596 in two different histories, pictured in Figure 6b and 6c, depending on the write with whom this  
 597 read event is associated by  $wr$ . The function VALIDWRITES limits the choices to those that preserve  
 598 consistency with the intended isolation level  $I$ , i.e.,  
 599

$$600 \quad 601 \quad \text{VALIDWRITES}(h, e) := \{w \mid h \oplus_j e \oplus wr(w, e) \text{ satisfies } I\}$$

## 603 5.2 Re-Ordering Events in Histories

604 After extending the current history with one more event, EXPLORE may recurse on other histories  
 605 obtained by re-ordering events in the current one (and dropping some other events).



618 Fig. 7. Example of inconsistency after swapping two events. TODO MAY NEED TO EXPLAIN THE CONVENTIONS IN THE PICTURE.

621

622

623 Re-ordering events must preserve the invariant of producing histories with at most one pending  
 624 transaction. To explain the use of this invariant in avoiding fruitless explorations, let us consider  
 625 the program in Figure 7a assuming an exploration under Read Committed. The oracle order gives  
 626 priority to the transaction on the left. Assume that the current history reached by the exploration  
 627 is the one pictured in Figure 7b (the last added event is write( $x, 2$ )). Swapping the write( $x, 2$ ) event  
 628 with the read( $x$ ) event would result in the history pictured in Figure 7c. To ensure that this swap  
 629 produces a new history which was not explored in the past, the  $wr_x$  dependency of read( $x$ ) is  
 630 changed towards the write( $x, 2$ ) transaction (we detail this later). By the definition of NEXT (and  
 631 the oracle order), this history shall be extended with read( $y$ ), and this read event will be associated  
 632 by  $wr_y$  to the only available write( $y, _$ ) event. This is pictured in Figure 7d. The next exploration  
 633 step will extend the history with write( $x, 2$ ) (the only extension possible) which however, results  
 634 in a history that does *not* satisfy Read Committed, thereby, the recursive exploration branch being  
 635 blocked. The core issue is related to the history in Figure 7d which has a pending transaction that is  
 636 *not* ( $so \cup wr$ )\*-maximal. Being able to extend such a transaction while maintaining consistency with  
 637

638 the intended isolation level is not guaranteed by Read Committed (and any other isolation level we  
 639 consider). Nevertheless, causal extensibility guarantees the existence of an extension for pending  
 640 transactions that are  $(\text{so} \cup \text{wr})^*$ -maximal. We enforce this requirement by restricting the explored  
 641 histories to have at most one pending transaction. This pending transaction will necessarily be  
 642  $(\text{so} \cup \text{wr})^*$ -maximal.

643  
 644 To enforce histories with at most one pending transaction, the function COMPUTEREORDERINGS,  
 645 which identifies events to reorder, has a non-empty return value only when the last added event is  
 646 commit (the end of a transaction). Therefore, in such a case, it returns pairs of read and write events  
 647 on the same variable, the write event coming from the last completed transaction, and such that  
 648 the transactions containing the two events are not causally dependent (i.e., related by  $[\text{so} \cup \text{wr}]^*$ ).  
 649

$$\begin{aligned} \text{COMPUTEREORDERINGS}(h_<) := & \{(r, w) \in \mathcal{E} \mid r < w \wedge r \in \text{reads}(t) \wedge w \in \text{writes}(t) \\ & \wedge \text{var}(r) = \text{var}(w) \\ & \wedge (\text{tr}(h, r), t) \notin [\text{so} \cup \text{wr}]^* \wedge t \text{ is complete}\} \end{aligned}$$

where  $t$  is the transaction log of the last event in  $<$

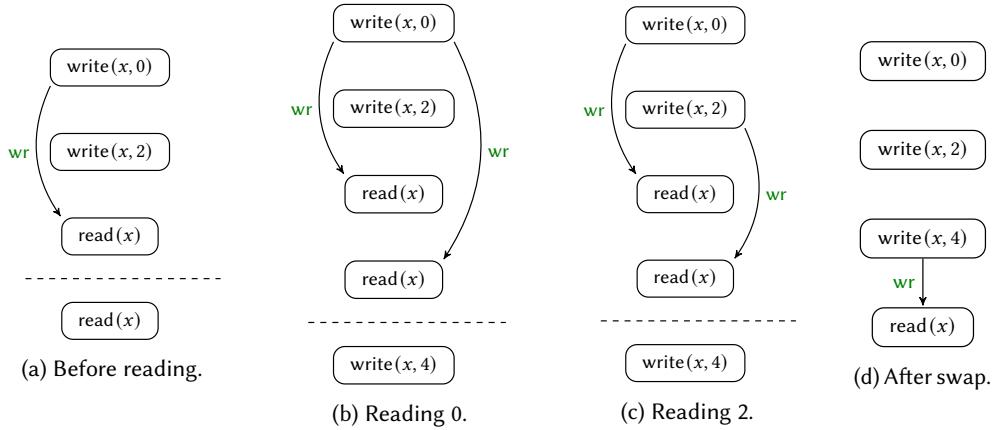
657 TODO GIVE AN EXAMPLE: A HISTORY, EMPHASIZING THE LAST COMPLETED TRANSACTION,  
 658 AND THE OUTPUT OF GENERIC COMPUTE. MAKE IT SO THE OUTPUT CONTAINS  
 659 TWO PAIRS. AND IT CAN BE USED TO EXEMPLIFY SWAP AS WELL (SOME THINGS TO KEEP  
 660 AS DEPENDENCIES, SOME THINGS TO DELETE, INCLUDING EVENTS THAT FOLLOW THE  
 661 READ IN THE SAME TRANSACTION.  
 662

663 Given a pair  $(r, w)$ , the function SWAP produces a new history  $h'$  which contains all events ordered  
 664 before  $r$  (w.r.t.  $<$ ), the transaction that contains  $w$  and all its  $[\text{so} \cup \text{wr}]^*$  predecessors, and the event  
 665  $r$  reading from  $w$ . Note that the  $\text{po}$  predecessors of  $r$  from the same transaction are ordered before  
 666  $r$  by  $<$  and they will be also included in  $h'$ . More generally, the history  $h'$  without  $r$  is a prefix of  
 667 the input history  $h$ . By definition, the only pending transaction in  $h'$  is the one containing the read  
 668  $r$ . The order relation is updated by moving the transaction containing the read  $r$  to be the last; it  
 669 remains unchanged for the rest of the events.  
 670

$$\begin{aligned} \text{SWAP}(h_<, r, w, \text{locals}) := & (h' = (h \setminus D) \oplus \text{wr}(w, r), <'), \text{locals}', \text{ where} \\ & D = \{e \mid r < e \wedge (\text{tr}(h, e), \text{tr}(h, w)) \notin [\text{so} \cup \text{wr}]^*\} \\ & <' = (< \downarrow (\text{events}(h') \setminus \text{events}(\text{tr}(h', r)))) \cdot \text{tr}(h', r) \\ & \text{locals}' = \text{locals} \downarrow \text{events}(h'). \end{aligned}$$

671  
 672 Above,  $h \setminus D$  denotes the prefix of  $h$  obtained by deleting all the events in  $D$  from its transaction  
 673 logs; a transaction log is removed all together if it becomes empty. Also,  $h'' \oplus \text{wr}(w, r)$  denotes  
 674 an update of the  $\text{wr}$  relation of  $h''$  where any pair  $(\_, r)$  is replaced by  $(w, r)$ . Finally,  $<'' \cdot \text{tr}(h', r)$   
 675 denotes an extension of the total order  $<''$  obtained by appending the events in  $\text{tr}(h', r)$  according  
 676 to program order.  
 677

678  
 679 TODO CONTINUE THE EXAMPLE ABOVE SHOWING THE RESULT OF SWAP IN ONE CASE.  
 680  
 681



701 Fig. 8. Re-ordering events versus optimality. Every history starts with a transaction writing the initial value 0.  
702  
703

### 704 5.3 Avoiding Redundancy

705 While extending histories according to NEXT and recursing on re-ordered histories whenever  
706 possible (taking OPTIMALITY as *true*) guarantees soundness and completeness, it does not guarantee  
707 optimality. There are two sources of redundancy in this trivial algorithm: (1) re-ordering the same  
708 read multiple times, and (2) applying SWAP on different histories may give the same result.  
709

710 TODO GIVE AN EXAMPLE FOR THE 1ST ISSUE. JUSTIFYING THE SWAPPED CONDITION.  
711

712 Without further restrictions on swaps, it is also possible that different branches of the recursion  
713 lead to the same history, thereby, violating optimality. As an example, consider a program with 2  
714 transactions that only read some variable  $x$  and 2 transactions that only write on  $x$ . Assume that  
715 EXPLORE reaches the ordered history pictured in Figure 8a and NEXT is about to return the second  
716 reading transaction. EXPLORE will recurse on the two histories pictured in Figure 8b and Figure 8c  
717 that differ in the write that this last read is reading from (either the initial write or the first write  
718 transaction). On both branches of the recursion, NEXT will extend the history with the last write  
719 transaction. For both histories, swapping this last write with the first read on  $x$  will result in the  
720 history pictured in Figure 8d (cf. the definition of COMPUTEORDERINGS and SWAP). Therefore,  
721 both branches of the recursion will continue extending the same history and optimality is violated.  
722 The source of non-optimality is related to *wr* dependencies that are *removed* when computing the  
723 result of SWAP. The histories in Figure 8b and Figure 8c were different because of the *wr* dependency  
724 involving the last read, but this difference was discarded during the computation of SWAP. Therefore,  
725 we will restrict the application of SWAP on histories where the discarded *wr* dependencies relate to  
726 some “fixed” set of writes, i.e., latest writes (w.r.t.  $<$ ) that are valid, i.e., preserve consistency with  
727 the intended isolation level (see the definition of  $\text{readsLatest}_I(\_, \_)$  below), e.g., when the second  
728 read reads from  $\text{write}(x, 2)$ .  
729

730 The OPTIMALITY condition, which restricts re-orderings, requires that every read that will be deleted  
731 by SWAP or the re-ordered read  $r$  (whose *wr* dependency will be modified) is not already swapped  
732 and it reads from a latest valid write ( $D$  has the same definition as in SWAP):  
733

**OPTIMALITY**( $h_<, r, w$ ) :=  $\forall r' \in \text{reads}(h) \cap (D \cup \{r\}). \neg \text{SWAPPED}(h_<, r') \wedge \text{readsLatest}_I(h_<, r')$

We say that a read  $r$  is *swapped* in  $h_<$  when (1) it reads from a write  $w$  that is a successor in the oracle order (the write was added by NEXT after the read), which is now a predecessor<sup>5</sup> in the history order  $<$  (2) there is no transaction  $t$  that is before  $r$  in both the oracle order  $<_{\text{or}}$  and the history order  $<$ , and which is a  $[\text{so} \cup \text{wr}]^+$  successor of  $w$ 's transaction, and (3)  $r$  is the first read in its transaction to read from  $w$ . Formally,

$$\begin{aligned}
& w < r \wedge w >_{\text{or}} r \\
& \wedge \\
\text{SWAPPED}(h_<, r) & := \forall e \in h. \text{tr}(h, e) <_{\text{or}} \text{tr}(h, r) \Rightarrow (r < e \vee (\text{tr}(h, w), \text{tr}(h, e)) \notin [\text{so} \cup \text{wr}]^+) \\
& \wedge \\
& \forall r' \in \text{reads}(h). (\text{tr}(h, w), r') \in \text{wr} \Rightarrow (r', r) \notin \text{po}
\end{aligned}$$

where  $w$  is the write event such that  $(w, r) \in \text{wr}$

A read  $r$  reads from a latest valid write, denoted as  $\text{readsLatest}_I(h_<, r)$ , if reading from any other later write w.r.t.  $<$  violates the isolation level  $I$ . Formally,

$\text{readsLatest}_I(h_<, r) := \forall w' \in \text{writes}(h). w < w' < r \Rightarrow (h \setminus \{e \mid r < e\}) \oplus \text{wr}(w', r) \not\models I)$   
 where  $w$  is the write event such that  $(w, r) \in \text{wr}$

## 5.4 Outputting histories

As we show in Section 8, the algorithm described above is  $I$ -sound assuming that  $\text{VALID}(h) := \text{true}$ . Therefore, any history computed by applying the functions described above will satisfy the intended isolation level  $I$ .

TODO THIS SHOULD PROBABLY BE ABOUT CORRECTNESS, STATING IT AND EXPLAINING A LITTLE BIT THE PROOFS

## 6 WEAK DPOR ALGORITHMS FOR SNAPSHOT ISOLATION AND SERIALIZABILITY

As shown in section 8, algorithm ??'s completeness proof (theorem 8.15) is model-depending, as it heavily relies on its causal-extensibility. Immediately, the question of the algorithm's extensibility to stricter isolation levels arise. For understanding the difference between the formers and SI or SER, let's analyze how algorithm ?? behaves for the program depicted in figure ?? under them. We study this example as it has exactly two consistent executions but a dead-lock one under SI and SER.

**THEOREM 6.1.** *If  $I$  is Snapshot Isolation or Serializability, there exists no EXPLORE algorithm that is  $I$ -sound,  $I$ -complete, and strongly optimal.*

<sup>5</sup>The EXPLORE maintains the invariant that every read follows the write it reads from in the history order <.

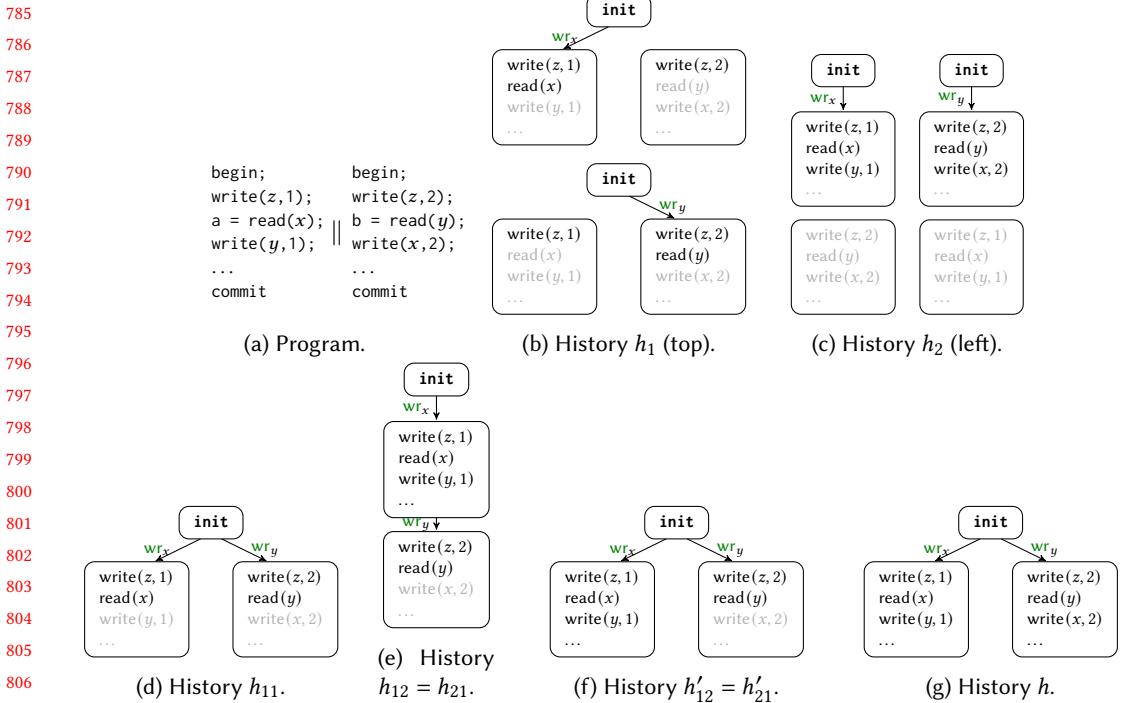


Fig. 9. A program and some partial histories. Events in grey are not yet added to the history. For  $h_{12}$  and  $h$ , the number of events that follow  $\text{write}(y, 1)$  and  $\text{write}(x, 2)$  is not important (we use black  $\dots$  to signify that).

PROOF. We consider the program in Figure 9a, and show that any concrete instance of the EXPLORE function in Algorithm 1 *can not be both I-complete and strongly optimal*. This program contains two transactions, where only the first 3 instructions in each transaction are important. We show that if EXPLORE is *I*-complete, then it will necessarily be called recursively on a history  $h$  like in Figure 9g which does not satisfy *I*, thereby violating strong optimality. In the history  $h$ , both *Snapshot Isolation* and *Serializability* forbid the two reads reading initial values while the writes following them are also executed (committed).

Assuming that the function NEXT is not itself blocking (which would violate strong optimality), the EXPLORE will be called recursively on *exactly one* of the four histories in Figure 9b and Figure 9c, depending on which of the two reads is returned first by NEXT and the order defined by NEXT between the writes. We will continue our discussion with the history  $h_1$  on the top of Figure 9b and the history  $h_2$  on the left of Figure 9c. The other cases are similar (symmetric).

From  $h_1$ , EXPLORE can be called recursively either on  $h_{11}$  in Figure 9d, or on  $h_{12}$  and  $h'_{12}$  in Figure 9e and Figure 9f, depending on the order defined by NEXT between  $\text{read}(y)$  and  $\text{write}(y, 1)$  ( $\text{read}(y)$  is returned by NEXT before  $\text{write}(y, 1)$  in  $h_{11}$  and vice-versa in  $h_{12}$  and  $h'_{12}$ ). The histories  $h_{12}$  and  $h'_{12}$  differ in the read-from associated to  $\text{read}(y)$ , and exploring at least  $h'_{12}$  is the best scenario towards ensuring *I*-completeness. If EXPLORE is called recursively only on  $h_{12}$ , then *I*-completeness is violated because  $h_{12}$  and any extension does not enable any re-ordering, and the history where  $\text{read}(x)$  reads from  $\text{write}(x, 2)$  will never be explored. Indeed, the two transactions in  $h_{12}$  are related by  $\text{wr}$  and events can be re-ordered earlier only together with their  $(\text{so} \cup \text{wr})^*$  predecessors. From

histories  $h_{11}$  and  $h'_{12}$ , EXPLORE will necessarily be called recursively on a history  $h$  like in Figure 9g which does not satisfy  $I$ , thereby violating strong optimality.

From  $h_2$ , EXPLORE can be called recursively on  $h_{21}$  in Figure 9e and  $h'_{21}$  in Figure 9f. As explained above for  $h_{12} = h_{21}$ , being called recursively only on  $h_{21}$  violates  $I$ -completeness, while being called recursively on  $h'_{21} = h'_{12}$  leads to an inconsistent history, thereby violating strong optimality.  $\square$

#### Present code

**THEOREM 6.2.** Algorithm **cite SER's algo** belongs to  $S_M \cap W_M$  for model  $M$ ,  $M \in \{\text{SI}, \text{SER}\}$ .

#### Discuss the exponential total branches ditched

## 7 THE PARTICULAR CASE OF PREFIX CONSISTENCY

*Prefix Consistency* is a special isolation level as it is not causal consistent **Example** but every partial history PRE-consistent can be extended into a consistent history 3.5. Therefore, neither algorithm **cite algorithm** is guaranteed to be sound, complete and strongly optimal for PRE nor the non-existence of such algorithm can be deduced in the same way as in theorem **cite SI and SER dont work**. In this section we present one program for which algorithm **cite algorithm** may or may not be consistent depending on the oracle order employed.

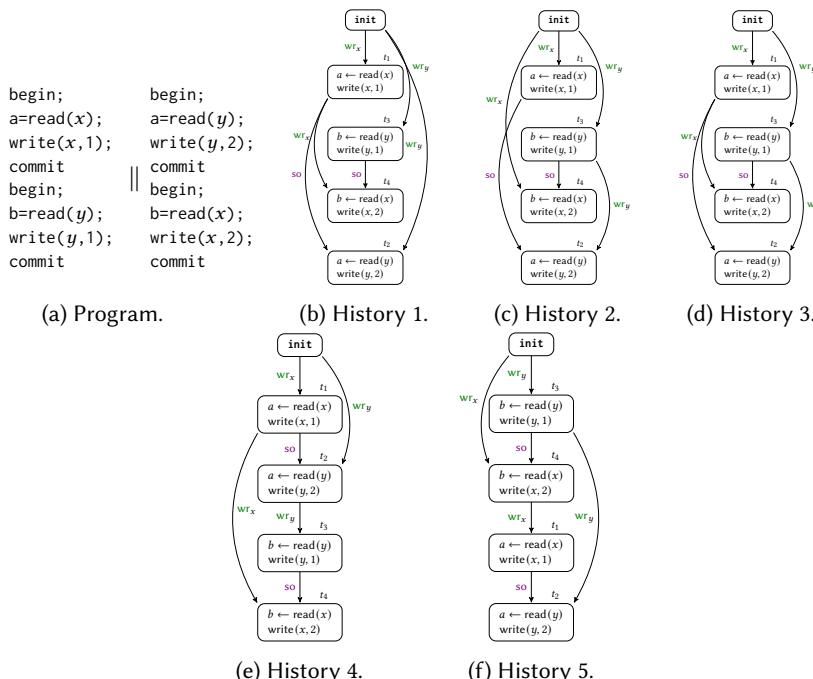


Fig. 10. A program along its PRE-consistent histories.

Let's analyze how algorithm **cite algorithm** would behave under program depicted in figure 10a; where  $<_{\text{or}}$  is defined as  $t_1 <_{\text{or}} t_2 <_{\text{or}} t_3 <_{\text{or}} t_4$ . In this situation, histories  $h_1$  and  $h_4$  (10b and 10e) are

883 computed without any call to SWAP function. It is clear that under  $h_4$ 's branch no swaps can be  
 884 produced as transactions are totally ordered via  $(\text{so} \cup \text{wr})^*$ . However, a swap between  $t_2$  and  $t_3$  can  
 885 be produced in the history  $h'_1 = h_1 \setminus t_4$  (which precedes  $h_1$ ) leading to histories  $h_2$  and  $h_3$  (figure  
 886 10c and 10d). However, as the read event in  $t_2$  in both  $h_2$  and  $h_3$  is swapped, we cannot swap  $t_4$   
 887 and  $t_1$  by  $\text{ISSWAPPABLE}_{\mathcal{M}}$ 's definition. Therefore,  $h_5$  (10f) is never reached.

888  
 889 On the other hand, let  $\text{or}'$  the oracle order defined as  $t_1 <_{\text{or}} t_3 <_{\text{or}} t_4 <_{\text{or}} t_2$ . In this case, histories  
 890  $h_1$ ,  $h_2$  and  $h_3$  can be obtained without any call to SWAP. Let  $h'_3 = h_3 \setminus t_2$ ; the maximum common  
 891 prefix of both  $h_1$  and  $h_3$ . Then, in  $h_3$  (respectively  $h'_3$ ), every event is maximally added, so  $t_2$  and  $t_3$   
 892 can be swapped to obtain  $h_4$  (respectively  $t_1$  and  $t_4$  to obtain  $h_5$ ). To conclude, with  $\text{or}'$  as oracle  
 893 order, the algorithm is complete.  
 894

## 8 PROOF OF THE ALGORITHMS

### Introduction

#### 8.1 Soundness

903 THEOREM 8.1. *Algorithm ?? is sound.*

904  
 905  
 906  
 907  
 908  
 909  
 910 PROOF. We prove this theorem by induction on the number of steps a reachable history needs in a  
 911 computable path to be reached. If this number is zero, the history is  $\emptyset$ , which is consistent; so let us  
 912 prove the inductive step, assuming that in any computable path of length at most  $n$  every history  
 913 is consistent. Let  $h$  computed in a  $n + 1$  path, and let  $h_p$  the immediate predecessor of  $h$ , which is  
 914 consistent, and  $a = \text{NEXT}(h_p)$ . If  $a$  is not a read event and  $h = h_p \bullet a$ , by NEXT's definition along  
 915 with the model's causal-extensibility,  $a$  is a  $\text{so} \cup \text{wr}$ -maximal event so  $h$  is consistent. Otherwise, if  
 916  $a$  is a read event and  $h = h_p \bullet_w a$  for some write event  $w$ , by VALID $_{\mathcal{M}}$ 's definition we know that  
 917  $h$  is consistent. Finally, if  $h = \text{SWAP}(h_p, r, w)$  for some events  $r, w$ , as  $\text{ISSWAPPABLE}_{\mathcal{M}}(h_p, r, w)$  is  
 918 satisfied,  $h$  is consistent.  $\square$

#### 8.2 Completeness

920  
 921  
 922  
 923  
 924  
 925 In our algorithm's context, completeness means being able to compute every total history. However,  
 926 our algorithm works with an extended version of histories where its events are totally ordered. For  
 927 proving this property, we will need to furnish every history with a total order that coincides with  
 928 the algorithm's one. This order is given by the canonical order function presented below.  
 929  
 930

---

**Algorithm 2 CANONICAL ORDER**


---

```

932 1: procedure CANONICALORDER( $h, T, T'$ )
933 2:   return  $T [so \cup wr]^* T' \vee$ 
934 3:    $(\neg(t(T')) [so \cup wr]^* T) \wedge \text{MINIMALDEPENDENCY}(h, T, T', \perp)$ 
935 4: end procedure

936 5: procedure MINIMALDEPENDENCY( $h, T, T', e$ )
937 6:   let  $a = \min_{<_{\text{or}}} \text{DEP}(h, T, e); a' = \min_{<_{\text{or}}} \text{DEP}(h, T', e)$ 
938 7:   if  $a \neq a'$  then
939 8:     return  $a <_{\text{or}} a'$ 
940 9:   else
941 10:    return MINIMALDEPENDENCY( $h, T, T', a$ )
942 11:   end if
943 12: end procedure

944 13: procedure DEP( $h, T, e$ )
945 14:   return  $\{r \mid \exists w \text{ s.t. } T [so \cup wr]^* t(w) \wedge w [wr] r \wedge t(r) [so \cup wr]^+ t(e)\} \cup T$ 
946 15: end procedure

```

---

951  
952 The function CANONICALORDER produces a relation between transactions in a history, denoted  $\leq^h$ .  
953 In algorithm 2's description, we denote  $\perp$  as the end of the program, which always exists, and that  
954 is **so**-related with every single transaction.  
955

956 LEMMA 8.2. *For every history  $h$ , event  $e$  and transaction  $T$ ,  $\text{DEP}(h, T, \min_{<_{\text{or}}} \text{DEP}(h, T, e)) \subseteq \text{DEP}(h, T, e)$ . Moreover, if  $\text{DEP}(h, T, e) \neq T$ , the inclusion is strict.*

957 PROOF. Let  $r' = \min_{<_{\text{or}}} \text{DEP}(h, T, e)$  and  $r \in \text{DEP}(h, T, r')$ . Then,  $\exists w \text{ s.t. } T [so \cup wr]^* t(w) \wedge w [wr] r \wedge t(r) [so \cup wr]^+ t(r')$  and  $\exists w' \text{ s.t. } T [so \cup wr]^* t(w') \wedge w' [wr] r' \wedge t(r') [so \cup wr]^+ t(e)$ ; so  
958  $t(r) [so \cup wr]^+ t(r') [so \cup wr]^+ t(e)$ . In other words,  $r \in \text{DEP}(h, T, e)$ . The moreover comes trivially  
959 as  $r' \notin \text{DEP}(h, T, r')$ .  $\square$

960 LEMMA 8.3. *For every distinct  $T, T'$ , MINIMALDEPENDENCY( $h, T, T', e$ ) always halts.*

961 PROOF. As  $h$  is a finite history, every transaction  $T$  belongs to  $\text{DEP}(h, T, e)$ , regardless of the event  
962  $e$  and via lemma 8.2 the set  $\text{DEP}$  shrinks in each recursive call; we conclude that if  $T \neq T'$ , there  
963 would be a call of MINIMALDEPENDENCY and an event  $e$  associated with it s.t.  $\min_{<_{\text{or}}} \text{DEP}(h, T, e) \neq \min_{<_{\text{or}}} \text{DEP}(h, T', e)$ .  $\square$

964 LEMMA 8.4. *The relation  $\leq^h$  is a total order.*

965 PROOF.

- 966 • Reflexivity: By definition, for every  $T$ ,  $T \leq^h T$ .
- 967 • Transitivity: Let's suppose  $a \leq^h b$  and  $b \leq^h c$ . First, take into account that if  $c \neq a$ ,  
968  $\neg(c [so \cup wr]^* a)$ . Here we distinguish four cases:

- 981 – If  $a [\text{so} \cup \text{wr}]^* b$  and  $b [\text{so} \cup \text{wr}]^* c$ , then  $a [\text{so} \cup \text{wr}]^* c$ , so  $a \leq^h c$ .
- 982 – If  $a [\text{so} \cup \text{wr}]^* b$  but  $\neg(b [\text{so} \cup \text{wr}]^* c)$ , then for every  $e \in h$ ,  $\min_{\text{or}} \text{DEP}(a, e) \leq_{\text{or}}$   
983  $\min_{\text{or}} \text{DEP}(b, e)$ , so  $a <^h c$ .
- 984 – If  $\neg(a [\text{so} \cup \text{wr}]^* b)$  but  $b [\text{so} \cup \text{wr}]^* c$ , then for every  $e \in h$ ,  $\min_{\text{or}} \text{DEP}(b, e) \leq_{\text{or}}$   
985  $\min_{\text{or}} \text{DEP}(c, e)$ , so  $a <^h c$ .
- 986 – If  $\neg(a [\text{so} \cup \text{wr}]^* b)$  and  $\neg(b [\text{so} \cup \text{wr}]^* c)$ , then it can be proven by induction that  
987  $a <^h c$ . It has to be proven iterating on the call function `mininalDependency`, a bit  
988 `boring`
- 989 • Antisymmetric For every  $a, b$  s.t.  $a \leq^h b$  and  $b \leq^h a$ . If  $a [\text{so} \cup \text{wr}]^* b$ , then  $a = b$ . If  
990 not, then `MINIMALDEPENDENCY`( $h, a, b, \perp$ ) and `MINIMALDEPENDENCY`( $h, b, a, \perp$ ) cannot be  
991 satisfied at the same time. Again an induction on `MINIMALDEPENDENCY` along with the  
992 history's finiteness.
- 993 • Strongly connection Let  $a, b$  s.t.  $a \not\leq_{\text{or}} b$ . If  $b [\text{so} \cup \text{wr}]^* a$ , then  $b \leq_{\text{or}} a$ . Otherwise, as  $\neg(a [\text{so} \cup$   
994  $\text{wr}]^* b)$  and `MINIMALDEPENDENCY` halts (lemma 8.3) and  $\neg\text{MINIMALDEPENDENCY}(h, a, b, e)$ ,  
995 then `MINIMALDEPENDENCY`( $h, b, a, e$ ); so  $b <^h a$ .
- 996 □

1000  
1001  
1002 *Definition 8.5.* A reachable history  $h$  is **or-respectful** if it has at most one pending transaction and  
1003 for every pair of events  $e \in \mathcal{P}, e' \in h$  s.t.  $e \leq_{\text{or}} e'$ , either  $e \leq_h e'$  or  $\exists e'' \in h, t(e'') \leq_{\text{or}} t(e)$  s.t.  
1004  $t(e') [\text{so} \cup \text{wr}]^* t(e'')$ ,  $e'' \leq_h e$  and `SWAPPED`( $h, e''$ ); where if  $e \notin h$  we state  $e' \leq_h e$  always hold  
1005 but  $e \leq_h e'$  never does. We will denote it by  $R^{\text{or}}(h)$ .

1006  
1007 LEMMA 8.6. Every reachable history is **or-respectful**.

1008  
1009 PROOF. We will prove it by induction on the number of ??'s stack calls a computable path that  
1010 leads to a history  $h$  needs,  $n$ . The base case,  $n = 0$ , is for the trivial history  $h = \emptyset$  where it trivially  
1011 holds; so let us prove the inductive case; being  $e \leftarrow \text{NEXT}(h)$ . On one hand,  $e$  is not a read nor a  
1012 begin event and  $h' = h \bullet e$ , as  $\neg\text{SWAPPED}(h, e)$  and  $h'$  is edge-wise identical to  $h$ ,  $R^{\text{or}}(h')$  holds.

1013  
1014 If  $e$  is a begin event,  $h' = h \bullet e$ . Let  $a \in \mathcal{P}, b \in h'$  s.t.  $a \leq_{\text{or}} b$ . If  $a \in h$  or  $b \neq e$ , as  $\leq_{h'}$  is an extension  
1015 of  $\leq_h$  and  $R^{\text{or}}(h)$ , the property holds. Moreover, as  $e = \min_{\text{or}} \mathcal{P} \setminus h$ , there is no event  $a \in \mathcal{P} \setminus h$  s.t.  
1016  $a \leq_{\text{or}} e$ ; so the property holds.

1017  
1018 On the other hand, if  $e$  is a read event and  $w$  is a write one, let us prove that  $h' = h \bullet_w e$ . Let  
1019  $a \in \mathcal{P}, b \in h'$  s.t.  $a \leq_{\text{or}} b$ . Once again, if  $a \in h$  or  $b \neq e$  the property holds; so let's suppose  
1020  $a \in \mathcal{P} \setminus h$  and  $b = e$ . Let  $d = \text{begin}(t(e))$ ,  $d \in h$ . As  $R^{\text{or}}(h)$  and  $a \notin h$ ,  $a \leq_{\text{or}} d$ ; so there exists  
1021  $c \in h, t(c) \leq_{\text{or}} t(a)$  s.t.  $t(d) [\text{so} \cup \text{wr}]^* t(c), c \leq_h d$  and `SWAPPED`( $h, c$ ). As  $t(r) = t(d)$ , we conclude  
1022  $R^{\text{or}}(h)$ .

1023 Finally, let  $h' = \text{SWAP}(h \bullet e, r, w)$  for some  $r, w \in h$  s.t. `ISSWAPPABLE_M`( $h \bullet e, r, w$ ) holds. Let  $a, b$   
1024 two event s.t.  $a \leq_{\text{or}} b$ . If  $a \leq_{h'} b$  or, as  $R^{\text{or}}(h)$  and `ISSWAPPABLE_M`( $h \bullet e, r, w$ ) holds,  $a \not\leq_h b$ , then  
1025 the property is satisfied; so let's suppose  $b <_{h'} a$  and  $a \leq_h b$ . In this situation,  $a$  has to be a deleted  
1026 event, so  $a \in \mathcal{P} \setminus h' \cup \{r\}$ . As  $r \leq_h a$ , if  $a \leq_{\text{or}} r$ , there would exist a  $c \in h, t(c) \leq_{\text{or}} t(a) \leq_{\text{or}} t(r)$  s.t.  
1027  $t(r) [\text{so} \cup \text{wr}]^* t(c)$  and `SWAPPED`( $h, c$ ). However, this contradicts `ISSWAPPABLE_M`( $h \bullet e, r, w$ ); so  
1028  $r \leq_{\text{or}} a$ . Taking  $e'' = r$  the property is witnessed. □

1030 PROPOSITION 8.7. For any reachable history  $h$ ,  $\leq^h \equiv \leq_h$ .

1031

1032

1033 PROOF. We will prove this lemma by induction on the number of steps a computable path leading  
 1034 to  $h$  are required by algorithm ???. The base case,  $n = 0$ , implies  $h = \emptyset$ , so both relations hold. Let's  
 1035 suppose that for every history  $h'$  that requires at most  $n$  steps,  $\leq^{h'} \equiv \leq_{h'}$ ; and let's analyze  $\leq^h$  for  
 1036 a history computed with  $n + 1$ . In particular, there exists a history  $h_p$  in that path which is an  
 1037 immediate predecessor of  $h$ . We will distinguish cases depending on how from  $h_p$  we reach  $h$ ;  
 1038 calling  $e = \text{NEXT}(h)$

1039 • Adding a end, write: As  $h_p$  and  $h$  are edge-wise identical,  $\leq^h \equiv \leq_h$ .

1040

1041 • Adding a begin: As  $\text{DEP}(h_p, T, \perp) = \text{DEP}(h, T, \perp)$  for every transaction in  $h_p$ , if  $T \leq^{h_p} T'$ ,  
 1042 then  $T \leq^h T'$ . Moreover,  $\text{DEP}(h, t(e), \perp) = \{e\} = \min_{\text{or}} \mathcal{P} \setminus h_p$ . By 8.6  $h$  is **or**-respectful, so  
 1043 for every  $T$ ,  $\min_{\text{or}} \text{DEP}(h, T, \perp) <_{\text{or}} e$ ; which implies  $T <^h t(e)$ . By lemma 8.4,  $\leq^h$  is a total  
 1044 order, so it coincides with  $\leq^h$ .

1045 • Adding a read: As no transaction depends on  $t(e)$  and  $t(e) = \text{last}(h_p)$ , if  
 1046 we prove that for every pair of transactions  $\text{MINIMALDEPENDENCY}(h_p, T, T', \perp) =$   
 1047  $\text{MINIMALDEPENDENCY}(h, T, T', \perp)$ , the lemma would hold. On one hand,  $\text{DEP}(h, t(e), \perp) =$   
 1048  $\text{DEP}(h_p, t(e), \perp) = t(e)$  and in the other hand, by lemma 8.6,  $\min_{\text{or}} \text{DEP}(h_p, T, \perp) <_{\text{or}} t(e)$ .  
 1049 Finally, as  $e \notin \text{DEP}(h, T, e')$ , for every  $T \neq t(e)$ ,  $e' \neq \perp$ , for every pair of transactions  $T, T'$ ,  
 1050  $\text{MINIMALDEPENDENCY}(h_p, T, T', \perp) = \text{MINIMALDEPENDENCY}(h, T, T', \perp)$ .

1051 • Swapping  $r \in h$  and  $w \in t(e)$ : As  $\text{ISSWAPPABLE}_{\mathcal{M}}(h \bullet e, r, w)$  is satisfied and  $h$   
 1052 is **or**-respectful, for every event  $e'$  and transaction  $T$ ,  $\min_{\text{or}} \text{DEP}(h_p, T, e') =$   
 1053  $\min_{\text{or}} \text{DEP}(h, T, e')$ , so for every pair of transactions  $\text{MINIMALDEPENDENCY}(h_p, T, T', \perp) =$   
 1054  $\text{MINIMALDEPENDENCY}(h, T, T', \perp)$ . In particular, this implies  $T \leq^{h_p} T'$  if and only if  $T \leq^h T'$   
 1055 for every pair  $T, T'$  and  $T \leq^h t(r)$ ; so  $\leq^h \equiv \leq_h$ .

1056

1057

1058

□

1059

1060

1061

1062 Proposition 8.7 is a very interesting result as it express the following fact: regardless of the com-  
 1063 putable path that leads to a history, the final order between events will be the same. This result  
 1064 will have a key role during both completeness and optimality, as it restricts the possible histories  
 1065 that precede another while describing the computable path leading to it. In addition, proposition  
 1066 8.7 together with lemma 8.6 justify enlarging definition 8.5 with the canonical order instead the  
 1067 computable order; and it is this new shape the one we will be using during the rest of proof.

1068

1069 LEMMA 8.8. Any total history is **or**-respectful.

1070

1071

1072 PROOF. Let  $h$  be a total history and  $T, T'$  a pair of transactions s.t.  $T \leq_{\text{or}} T'$ . If  $T \leq^h T'$ , then the  
 1073 statement is satisfied; so let's assume the contrary:  $T' \leq^h T$ . If  $T' [\text{so} \cup \text{wr}]^* T$ , then for every  
 1074  $e \in T, e' \in T' \exists c \in h$  s.t.  $t(c) \leq_{\text{or}} t(e), t(e') [\text{so} \cup \text{wr}]^* t(c)$ ,  $\text{SWAPPED}(h, c)$  and  $c \leq^h e$ ; so the  
 1075 property is satisfied. Otherwise, by definition of **MINIMALDEPENDENCY**, there exists  $r' \in h$  s.t.  
 1076  $T' [\text{so} \cup \text{wr}]^* t(r')$  and  $t(r') \leq_{\text{or}} T$ . Moreover, by **CANONICALORDER**'s definition,  $t(r) \leq^h T$ . Finally  
 1077  $\text{SWAPPED}(h, r')$  holds as it is the minimum element according **or**. To sum up,  $R^{\text{or}}(h)$  holds. □

1078

---

**Algorithm 3** PREV
 

---

```

1079
1080
1081 1: procedure PREV( $h$ )
1082   2:   if  $h = \emptyset$  then
1083     3:     return  $\emptyset$ 
1084   4:   end if
1085   5:    $a \leftarrow \text{last}(h)$ 
1086   6:   if  $\neg\text{SWAPPED}(h, a)$  then
1087     7:     return  $h \setminus a$ 
1088   8:   else
1089     9:     return MAXCOMPLETION( $h \setminus a, \{e \mid e \notin (h \setminus a) \wedge e <_{\text{or}} h.\text{wr}(a)\}$ )
1090   10:   end if
1091 11: end procedure

1092
1093 12: procedure MAXCOMPLETION( $h, D$ )
1094   13:   if  $D \neq \emptyset$  then
1095     14:      $e \leftarrow \min_{<_{\text{or}}} D$ 
1096     15:     if  $e.\text{type}() \neq \text{read}$  then
1097       16:       return MAXCOMPLETION( $h \bullet e, D \setminus \{e\}$ )
1098     17:     else
1099       18:       let  $w$  s.t.  $\text{readsLatest}_I(h \bullet_w e, e)$ 
1100       19:       return MAXCOMPLETION( $h \bullet_w e, D \setminus \{e\}$ )
1101   20:   end if
1102   21:   else
1103     22:     return  $h$ 
1104   23:   end if
1105 24: end procedure

```

---

1106  
 1107 Function 3 produce a history that are meant to be the previous step of a reachable history. Thanks  
 1108 to this definition, we will show that every total history has a computable path based on applying  
 1109  $\text{PREV}^{-1}$  function iteratively until the objective history is reached.

1110  
 1111 TODO (somewhere before): if  $h \rightarrow \text{SWAP}(h \bullet e, r, w)$  “in one step”, actually from  $h$  we go to  $h \bullet e$   
 1112 and from it to the swapped.

1113  
 1114 LEMMA 8.9. For every or-respectful history  $h$ ,  $\text{PREV}(h)$  is also or-respectful.

1115  
 1116 PROOF. Let suppose  $h \neq \emptyset$ ,  $h_p = \text{PREV}(h)$ ,  $a = \text{last}(h)$ ,  $e \in \mathcal{P}$  and  $e' \in h_p$  s.t.  $e \leq_{\text{or}} e'$ . As  
 1117  $R^{\text{or}}(h)$  is satisfied, either  $e \leq^h e'$  or  $\exists e'' \in h, t(e'') \leq_{\text{or}} t(e), e'' \leq^h e, t(e'')$  [ $\text{so} \cup \text{wr}]^* t(e'')$  and  
 1118  $\text{SWAPPED}(h, e'')$ . If  $\neg\text{SWAPPED}(h, a)$ ,  $h_p = h \bullet a$ ; so if  $e \leq^h e'$ ,  $e \leq^{h_p} e'$  and if not,  $e'' \in h_p$ , so  $R^{\text{or}}(h_p)$   
 1119 holds.

1120  
 1121 Otherwise,  $\text{SWAPPED}(h, a)$  and we distinguish between the sets  $e$  and  $e'$  belong to. Firstly, for  
 1122 every pair of events  $\hat{e} \in h_p \setminus h$ ,  $\hat{e}' \in \text{DEP}(h, t(\hat{e}, \perp))$ , we know that  $t(\hat{e}) \leq_{\text{or}} t(\hat{e}')$ . Therefore,  
 1123  $\min_{<_{\text{or}}} \text{DEP}(h, t(\hat{e}, \perp)) = \text{begin}(t(\hat{e}))$ . In addition, by construction of  $\text{PREV}(h)$  and or-respectfulness  
 1124 of  $h$ , for every  $\hat{e} \in h$ ,  $e'' \in h$ ,  $\min_{<_{\text{or}}} \text{DEP}(h_p, t(\hat{e}), e'') = \min_{<_{\text{or}}} \text{DEP}(h, t(\hat{e}), e'')$ . Combining both  
 1125 results, if  $e'$  belong to  $h$ , either  $e \leq^{h_p} e'$  or exists a  $e'' \in h$  s.t.  $e'' \leq^{h_p} e$  and witness  $R^{\text{or}}(h)$  for  $e, e'$   
 1126 (regardless of  $e$ 's belonging to  $h$ ,  $e'' \leq^{h_p} e$ ). On the contrary, as  $h_p$  has no pending transactions, if  
 1127

1128    $e' \notin h$ ,  $\neg(t(e'))$  [**so**  $\cup$  **wr**] $^*$   $t(e)$ ), so regardless if  $t(e)$  [**so**  $\cup$  **wr**] $^*$  $t(e')$ ,  $e \leq^{hp} e'$ . To sum up,  $R^{\text{or}}(h_p)$   
 1129   holds.  $\square$

1130

LEMMA 8.10. For every consistent history or-respectful  $h$ , if  $\text{PREV}(h)$  is reachable, then  $h$  is also reachable.

1133

PROOF. Let suppose  $h \neq \emptyset$ ,  $h_p = \text{PREV}(h)$  and  $a = \text{last}(h)$ . If  $\neg\text{SWAPPED}(h, a)$ , let  $h_n = h_p \bullet a$  if  $a$  is not a read,  $h_n = h_p \bullet_{h.\text{wr}(a)} a$  in the other case. Either way,  $h_n$  is always reachable and it coincides with  $h$ . Otherwise,  $a$  is a read event and it swapped; so let us call  $w = h.\text{wr}(a)$ . Firstly, as  $\text{SWAPPED}(h, a)$ ,  $a <_{\text{or}} w$ , and by lemma 8.6,  $R^{\text{or}}(h_p)$  holds, so  $a <_{h_p} w$  does; which let us conclude  $\text{COMPUTE}(h_p)$  will always return  $(a, w)$  as a possible swap pair. In addition, all transactions in  $h_p$  are non-pending, so in particular  $\text{last}(h_p)$  is an end event. If we call  $h_s = \text{SWAP}(h_p, a, w)$ , and  $h_p \setminus h = h_p \setminus h_s$  would hold, as  $h \subseteq h_p, h_s \subseteq h_p$ , then  $h = h_s$ ; which would allow us to conclude  $h$  is reachable from  $h_p$ .

On one hand, if  $e \in h_p \setminus h$ ,  $e \notin h$  and  $e <_{\text{or}} w$ . In particular,  $\neg(t(e) [\text{so} \cup \text{wr}]^* t(w))$ . Moreover, if  $e \leq_{\text{or}} a$ , by  $R^{\text{or}}(h)$ , either  $e \leq^h a$  or  $\exists e'' \in h, e'' \leq_{\text{or}} e$  s.t.  $t(a) [\text{so} \cup \text{wr}]^* t(e'')$ ,  $e'' \leq^h e$  and  $\text{SWAPPED}(h, e'')$ ; both impossible situations as  $e \notin h$  and  $a = \text{last}(h)$ ; so  $a \leq_{\text{or}} e$ . In other words,  $e \in h_p \setminus h_s$ .

On the other hand,  $e \in h_p \setminus h_s$  if and only if  $\neg(t(e) [\text{so} \cup \text{wr}]^* t(w))$  and  $a <_{\text{or}} e <_{\text{or}} w$ . If  $e \in h$  then  $e \leq^h a$ , and as  $h$  is **or**-respectful and  $a \leq_{\text{or}} e$ , we deduce there exists a  $e'' \in h$  s.t.  $t(e'') \leq_{\text{or}} t(a)$ ,  $t(e) [\text{so} \cup \text{wr}]^* t(e'')$  and  $\text{SWAPPED}(h, e'')$ . Moreover, as  $c \in h$ ,  $c \in h_p$ ; but as  $\text{SWAPPED}(h_p, c)$  and  $\text{ISSWAPPABLE}_{\mathcal{M}}(h, a, w)$  hold,  $c \in h_s$  and so  $e$  does. This result leads to a contradiction, so  $e \notin h$ ; i.e.  $e \in h_p \setminus h$ .  $\square$

**COROLLARY 8.11.** In a consistent or-respectful history  $h$  whose previous history is reachable, if its last event  $a$  is swapped,  $h$  coincides with  $\text{SWAP}(\text{PREV}(h), a, h.\text{wr}(a))$ .

**PROOF.** It comes straight away from the proof of lemma 8.10. □

**LEMMA 8.12.** *For every non-empty consistent or-respectful history  $h$ ,  $h_p = \text{PREV}(h)$  and  $a = \text{last}(h)$ , if  $\text{SWAPPED}(h, a)$  then  $\{e \in h_p \mid \text{SWAPPED}(h_p, e)\} = \{e \in h \mid \text{SWAPPED}(h, e)\} \setminus \{a\}$ , otherwise  $h_p = h \setminus a$ .*

**PROOF.** Let  $a = \text{last}(h)$  and  $h' = h \setminus a$ . If  $a$  is not swapped, then  $h_p = h'$ , so the lemma holds immediately. Otherwise, as  $h_p = \text{MAXCOMPLETION}(h')$ , we will show that every event not belonging to  $h_p \setminus h'$  is not swapped by induction on every recursive call to `MAXCOMPLETION`. Let us call  $D = \{e \mid e \notin h' \wedge e <_{\text{or}}\}$ . This set, intuitively, contain all the events that would have been deleted from a reachable history  $h$  to produce  $h_p$ . In this setting, let us call  $h|_{|D|} = h'$ ,  $D_{|D|} = D$  and  $D_k = D_{k+1} \setminus \{\min_{<_{\text{or}}} D_{k+1}\}$ ,  $e_k = \min_{<_{\text{or}}} D_k$  for every  $k, 0 \leq k < |D|$  (i.e.  $D_k = D_{k+1} \setminus \{e_{k+1}\}$ ). We will prove the lemma by induction on  $n = |D| - k$ , constructing a collection of histories  $h_k$ ,  $0 \leq k < |D|$ , such that each one is an extension of its predecessor with a non-swapped event.

The base case,  $h|_{D'}$  is trivial as by its definition it corresponds with  $h'$ . Let's prove the inductive case:  $\{e \mid \text{SWAPPED}(h_{k+1}, e)\} = \{e \mid \text{SWAPPED}(h', e)\}$ . If  $e_{k+1}$  is not a read event,  $h_k = h_{k+1} \bullet e_{k+1}$  and  $\{e \mid \text{SWAPPED}(h_k, e)\} = \{e \mid \text{SWAPPED}(h', e)\}$ ; as only read events can be swapped. Otherwise, by the model's causal-extensibility there exists a write event  $f_{k+1}$  s.t. writes the same variable and  $\text{isConsistent}_{\mathcal{M}}(h_{k+1} \bullet_{f_{k+1}} e_{k+1}) \wedge t(f_{k+1}) [\text{so} \cup \text{wr}]^* t(e_{k+1})$  holds.  $\{e \mid \text{SWAPPED}(h_{k+1}, e)\} = \{e \mid \text{SWAPPED}(h_{k+1} \bullet_{f_{k+1}} e_{k+1}, e)\}$  holds. Let  $E_{k+1} = \{w \mid \text{isConsistent}_{\mathcal{M}}(h_{k+1} \bullet_w e_{k+1}) \wedge$

1177  $\{e \mid \text{SWAPPED}(h_{k+1}, e)\} = \{e \mid \text{SWAPPED}(h_{k+1} \bullet_w e_{k+1}, e)\}$  and  $w_{k+1} = \max_{\leq h_{k+1}} E_{k+1}$ . This el-  
 1178 ement is well defined as  $f_{k+1}$  belongs to  $E_{k+1}$ . Therefore,  $h_k = h_{k+1} \bullet_{w_{k+1}} e_{k+1}$  is consistent  
 1179 and  $\{e \mid \text{SWAPPED}(h_k, e)\} = \{e \mid \text{SWAPPED}(h', e)\}$ . Moreover, let's remark that as  $w_{k+1}$  is the  
 1180 maximum write event according to  $\leq_{h_{k+1}}$  s.t.  $\text{isConsistent}_{\mathcal{M}}(h_k)$  and  $\{e \mid \text{SWAPPED}(h_k, e)\} =$   
 1181  $\{e \mid \text{SWAPPED}(h', e)\}$  and  $R^{\text{or}}(h)$ , it also satisfies  $\text{readsLatest}_I(h_k, e_{k+1}, w_{k+1})$ . Altogether, we obtain  
 1182  $h_p = h_0$ ; which let us conclude  $\{e \in h_p \mid \text{SWAPPED}(h_p, e)\} = \{e \in h' \mid \text{SWAPPED}(h', e)\} = \{e \in$   
 1183  $h \mid \text{SWAPPED}(h, e)\} \setminus \{a\}$ .  $\square$

1184  
 1185 LEMMA 8.13. *For every history  $h$  there exists some  $k_h \in \mathbb{N}$  such that  $\text{PREV}^{k_h}(h) = \emptyset$ .*

1186  
 1187 PROOF. This lemma is immediate consequence of lemma 8.12. Let us call  $\xi(h) = |\{e \in$   
 1188  $h \mid \text{SWAPPED}(h, e)\}$ , the number of swapped events in  $h$ , and let us prove the lemma by induction  
 1189 on  $(\xi(h), |h|)$ . The base case,  $\xi(h) = |h| = 0$  is trivial as  $h$  would be  $\emptyset$ ; so let's assume that  
 1190 for every history  $h$  such that  $\xi(h) < n$  or  $\xi(h) = h \wedge |h| < m$  there exists such  $k_h$ . Let  $h$  then a  
 1191 history s.t.  $\xi(h) = n$  and  $|h| = m$ .  $h_p = \text{PREV}(h)$ . On one hand, if  $h_p = h \setminus a$  then  $\xi(x_p) = \xi(h)$   
 1192 and  $|h_p| = |h| - 1$ . On the other hand, if  $h_p \neq h \setminus a$ ,  $\xi(h_p) = \xi(h) - 1$ . In any case, by induction  
 1193 hypothesis on  $h_p$ , there exists an integer  $k_{h_p}$  such that  $\text{PREV}^{k_{h_p}}(h_p) = \emptyset$ . Therefore,  $k_h = k_{h_p} + 1$   
 1194 satisfies  $\text{PREV}^{k_h}(h) = \emptyset$ .  $\square$

1195  
 1196 PROPOSITION 8.14. *For every consistent or-respectful history  $h$  exists  $k \in \mathbb{N}$  and some sequence of*  
 1197 *or-respectful histories  $\{h_n\}_{n=0}^k$ ,  $h_0 = \emptyset$  and  $h_k = h$  such that the algorithm will compute.*

1198  
 1199 PROOF. Let  $h$  a history,  $k$  the minimum integer such that  $\text{PREV}^k(h) = \emptyset$ , which exists thanks to  
 1200 lemma 8.13 and  $C = \{\text{PREV}^{k-n}(h)\}_{n=0}^k$  a set of indexed histories. By the collection's definition and  
 1201 lemma 8.9,  $h_0 = \text{PREV}^k(h) = \emptyset$ ,  $h_k = \text{PREV}^0(h) = h$  and  $R^{\text{or}}(h_n)$  for every  $n \in \mathbb{N}$ ; so let us prove by  
 1202 induction on  $n$  that every history in  $C$  is reachable. The base case,  $h_0$ , is trivially achieved; as it is  
 1203 always reachable. In addition, by lemma 8.10, we know that if  $h_n$  is reachable,  $h_{n+1}$  is it too; which  
 1204 proves the inductive step.  $\square$

1205  
 1206 THEOREM 8.15. *Algorithm ?? is complete.*

1207  
 1208 PROOF. By lemma 8.8, any consistent total history is or-respectful. As a consequence of proposition  
 1209 8.14, there exist a sequence of reachable histories which  $h$  belongs to; so in particular,  $h$  is reachable.  
 1210  $\square$

1211  
 1212 **8.3 Optimality**

1213 I have this new proof shorter than the initial idea but I do not know if it doesn't have typos. TODO:  
 1214 reread it!

1215  
 1216  
 1217 LEMMA 8.16. *For every two histories  $h_1, h_2$  with the same set of events and read event  $r$ , if*  
 1218  *$\text{ISWAPPABLE}_{\mathcal{M}}(h_i, r, w)$  holds for  $i \in \{1, 2\}$  and  $\leq_{h_i} \equiv \leq_{h_2}$ , for every read event  $r'$  s.t.  $r \leq_{h_i} r'$*   
 1219 *and  $\neg(t(r') [\text{so} \cup \text{wr}]^+ t(w))$  we have  $h_1.\text{wr}(r') = h_2.\text{wr}(r')$ .*

1220  
 1221 PROOF. As for every  $i \in \{1, 2\}$   $\text{ISWAPPABLE}_{\mathcal{M}}(h_i, r, w)$  holds, so does it  $\text{readsLatest}_I(h_i, r')$ . There-  
 1222 fore, as  $\leq_{h_1} \equiv \leq_{h_2}$ ,  $h_1.\text{wr}(r') = h_2.\text{wr}(r')$ .  $\square$

1223  
 1224 THEOREM 8.17. *Algorithm ?? is optimal.*

1226 PROOF. As the model is causal-extensible, any algorithm weakly optimal is also optimal. Let us  
 1227 prove that for every reachable history there is only a computable path that leads to it from  $\emptyset$ . By  
 1228 lemma 8.7, we know that for every reachable history  $h$ ,  $\leq_h \equiv \leq^h$ . However,  $\leq^h$  is an order that does  
 1229 not depend on the computable path that leads to  $h$ ; so neither does  $\leq_h$ . Let's suppose that there  
 1230 exist two reachable histories  $h_1, h_2$  s.t.  $h_1 = \text{PREV}(h)$  and that in one step of computation produce  
 1231 a common history  $h$ . If we prove they are identical, the algorithm would be optimal. Firstly, if  
 1232  $\text{last}(h)$  is not a swapped read event, by the definition of NEXT function  $h_2 = h \setminus \text{last}(h) = h_1$ .  
 1233 On the contrary, let's suppose  $r = \text{last}(h)$  is a swapped event that reads from a write event  
 1234  $w$ . Because  $\text{SWAPPED}(h, r)$  holds, from  $h_2$  to  $h$  it has to have happened a swap between  $r$  and  $w$ .  
 1235 But by corollary 8.11,  $h = \text{SWAP}(h_1, r, w)$ , so  $h_1 \upharpoonright_{h \setminus r} = h_2 \upharpoonright_{h \setminus r}$ . As  $h_1, h_2$  are both or-respectful,  
 1236 if  $e \in h_1 \setminus h \iff e \in h_2 \setminus h$ . Finally, as  $\text{ISWAPPABLE}_{\mathcal{M}}(h_i, r, w)$  holds for  $i \in \{1, 2\}$ , by lemma 8.16,  
 1237 if  $e \in h_1 \setminus h$ ,  $h_1.\text{wr}(e) = h_2.\text{wr}(e)$ . Therefore,  $h_1 = h_2$ .  $\square$

1238

1239

1240

1241

1242

1243

1244

1245

1246

1247

1248

1249

1250

1251

1252

1253

1254

1255

1256

1257

1258

1259

1260

1261

1262

1263

1264

1265

1266

1267

1268

1269

1270

1271

1272

1273

1274